A free-space convolution on a staggered mesh for the initial field of a uniformly moving particle bunch

René Widera
Helmholtz-Zentrum Dresden-Rossendorf (HZDR), Dresden, Germany
r.widera@hzdr.de

Abstract

Particle-in-cell simulations of laser wakefield acceleration need the initial electromagnetic field of a relativistic bunch before the first push, and the current workaround, a few field-free steps from a zero-field state, has no physical basis. We compute the present-time field of a uniformly moving bunch as a zero-padded free-space FFT convolution of a TSC charge deposit with the sampled TSC-convolved Liénard–Wiechert kernel, read out on the staggered Yee grid, and we measure its accuracy and cost. Against an independent quadrature reference the mesh reaches 3×10^{-4} L2-pooled relative error within eight cells of the bunch axis for a stationary particle and within forty-eight cells for $\gamma = 10$ perpendicular to the beam. The deposit dominates the compute time at large particle counts, growing with N and approaching linear in N at the top of the sweep, while the FFT scales with the padded grid; the total for 2×10^6 particles in a 20-cell box is 51.9 ms on one H100 of JURECA H100 (compute-only, median of nine repeats). An opt-in near shell replaces the mesh inside the TSC support with a closed-form evaluator, assembles exactly to roundoff, and is bitwise independent of the MPI rank count. The method writes a float32 openPMD checkpoint that restarts under a stock PIConGPU.

1 Introduction

When a particle-in-cell (PIC) simulation of laser wakefield acceleration starts, the driven bunch already moves relativistically through the plasma [Esarey et al., 2009], so the first push must act on the correct electromagnetic force. In practice the start state carries no field: simulation frameworks either start field-free and let the wake build up in the first few steps, or compute the initial field by evaluating the closed-form expression of each source charge at every observation point. The first option inserts a transient that the physics of a uniformly moving bunch does not contain, and the second is too slow once the staggered mesh holds thousands of observation points per source.

The present-time field of a uniformly moving charge is known in closed form [Jackson, 2007, Landau and Lifshitz, 1975], and the field of a bunch of N such charges is the linear sum over sources. Evaluated on a mesh that also carries a TSC deposit [Birdsall and Langdon, 1985, Hockney and Eastwood, 1981], the sum becomes a discrete convolution between the deposited charge and the TSC-convolved point-charge kernel. A zero-padded free-space FFT evaluates that convolution in $\mathrm{O}(M^3 \log M)$ transforms with M the padded grid size, independent of N in the transform itself. `picBunch` implements exactly this operation, with staggered Yee read-out, an opt-in near shell that evaluates the in-support integral in closed form, an MPI-parallel component- and tile-parallel design, and a checkpoint writer that restarts under a stock PIConGPU [Burau et al., 2010, Huebl et al., 2018].

This paper gives the discrete recipe, its physical consistency checks, and measured accuracy and cost. Section 2 defines the field and Section 3 the algorithm. Section 4 checks the kernel against the Liénard–Wiechert solution, Lorentz contraction, Gauss's law, the Yee staggering, and the $B = \boldsymbol{v} \times \boldsymbol{E}/c^2$ limit. Section 5 separates the error into mesh/deposit, velocity binning, near-source quadrature, far-tail, and output-precision terms, and states the domain of validity of each near-shell mode. Section 7 measures cost and MPI scaling. Appendices give the full derivations, the secondary validation sweeps, the large-domain slices, and the multi-velocity measurements.

2 Physical model and field definition

The computed quantity is a single linear integral. For a bunch of source charges at positions $\mathbf{x}_p$ with signed charge q_p and a shared dimensionless velocity $\mathbf{u} = \mathbf{p}/(mc)$, the field at a staggered observation point $\mathbf{x}_{\text{obs}}$ is the TSC-convolved uniformly-moving point-charge field

$$F_c(\mathbf{x}_{\text{obs}}) = \sum_p q_p \int_{[-1.5,1.5]^3} W(\mathbf{s})\, K_{\text{point},c}\big(\mathbf{x}_{\text{obs}} - \mathbf{x}_p - \mathbf{s}\,\mathbf{cell},\ \mathbf{u}\big)\, d\mathbf{s}^3, \tag{1}$$

where $W(\mathbf{s}) = W_x(s_x)W_y(s_y)W_z(s_z)$ is the separable TSC shape (Section 2.3) and K_{point} is the per-unit-charge field of one uniformly moving charge. The sum over particles is exact because the integral is linear in the source charge; the algorithm of Section 3 is a fast evaluation of this one integral.

2.1 The moving point kernel

The kernel is the present-time Liénard–Wiechert field of a uniformly moving charge in the present-position form [Jackson, 2007, Landau and Lifshitz, 1975]. In the lab frame, with $\mathbf{R} = \mathbf{x}_{\text{obs}} - \mathbf{x}_{\text{src}}$,

$$\gamma = \sqrt{1 + |\mathbf{u}|^2}, \qquad \boldsymbol{\beta} = \mathbf{u}/\gamma, \qquad \text{invg}^2 = 1/(1 + |\mathbf{u}|^2) = 1 - \beta^2, \tag{2}$$

$$D = |\mathbf{R}|^2\, \text{invg}^2 + (\mathbf{R}\cdot\boldsymbol{\beta})^2, \tag{3}$$

$$\mathbf{E}_n = \text{invg}^2\, \mathbf{R}/D^{3/2}, \qquad \mathbf{B}_n = \boldsymbol{\beta} \times \mathbf{E}_n. \tag{4}$$

For $\mathbf{u} = \mathbf{0}$ exactly, $\mathbf{E}_n = \mathbf{R}/|\mathbf{R}|^3$ and $\mathbf{B}_n = \mathbf{0}$; no ε is added to any denominator. The kernel is homogeneous of degree -2 in $\mathbf{R}$. A bunch with a velocity spread needs one convolution per velocity class (Section D.1, Appendix D); the validated scope is a single shared $\mathbf{u}$.

2.2 The unit chain

The kernel normalizes the per-unit-charge field as $E_n = 4\pi\varepsilon_0\, E/q$ and $B_n = 4\pi\varepsilon_0 c\, B/q$, so $B_n = \boldsymbol{\beta} \times E_n$. The mesh produces the per-unit-charge field with $\mathbf{R}$ in metres, already summed over the deposited charge q_p in coulombs. The chain to physical SI is

$$E_{\text{SI}}[c] = F_c[c]\, h^3/(4\pi\varepsilon_0), \quad B_{\text{SI}}[c] = F_c[c]\, h^3/(4\pi\varepsilon_0 c), \tag{5}$$

with $\varepsilon_0 = 8.8541878128 \times 10^{-12}\,\text{F/m}$ and $c = 299792458\,\text{m/s}$. The h^3 factor is the Riemann volume of the sampled continuum integral (Section 3).

2.3 The TSC shape

W is the quadratic B-spline, with argument in cell units, support $[-1.5, 1.5]$ per axis, and unit integral per axis:

$$W(x) = \begin{cases} 3/4 - x^2 & |x| < 1/2, \\ \frac{1}{2}(3/2 - |x|)^2 & 1/2 \le |x| < 3/2, \\ 0 & \text{otherwise.} \end{cases} \tag{6}$$

3 Numerical algorithm

The mesh evaluates Eq. (1) by a discrete deposit, a sampled kernel, a free-space FFT convolution, and a staggered read-out. Let N be the bunch-box nodes per axis, $M = \text{pad} \cdot N$ the padded FFT grid per axis, h the node spacing in cell units (node i sits at $ih\,\mathbf{cell}$), and $\mathbf{cell} = (dx, dy, dz)$. All precision-critical work is in double precision.

3.1 Field evaluation: mesh plus near correction

The field of Section 2 is assembled as a smooth part plus a near-field correction (Fig. 1). The smooth part, F_{mesh}, is the zero-padded free-space FFT convolution of the TSC deposit with the sampled TSC-convolved kernel (Sections 3–3.5); it is the far field and is what is written when no correction is requested. A near correction removes the mesh's near-field error, where the kernel is singular. The total is

$$F = F_{\text{mesh}}(\text{all}) - F_{\text{mesh}}(\text{near}) + F_{\text{near}}(\text{near}), \tag{7}$$

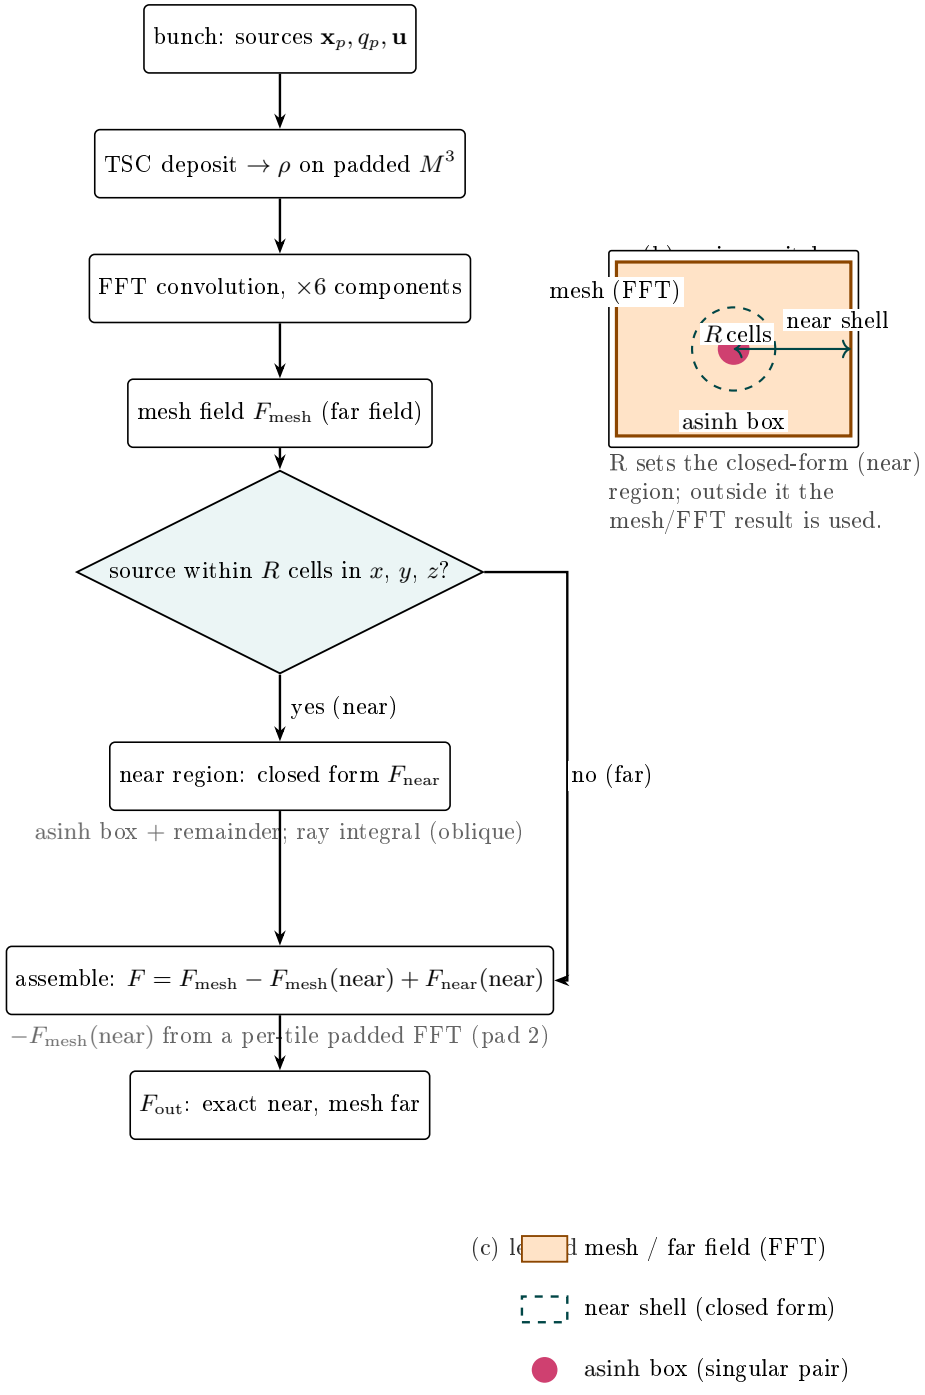


Figure 1: How the field is computed. **(a)** The pipeline: a TSC deposit and a free-space FFT convolution produce the mesh field F_{mesh}, accurate in the far field. The switch asks, per (source, node) pair, whether the pair is within the near-shell radius R *on every axis*; the far pair keeps F_{mesh}, the near pair is replaced by the closed-form pair integral (an asinh box plus a smooth remainder, a ray integral for oblique motion). Assembly subtracts the mesh value of the near pairs (a per-tile zero-padded FFT, padding 2) and adds the closed form, so the output is exact in the near region and mesh-accurate in the far region. **(b)** The region switch in real space: the mesh (orange) holds everywhere except a per-axis box of half-width R about each source (teal ring), and the asinh box acts where the observer is inside the TSC support (purple core, 1.5 cells).

where F_{near} is the closed-form TSC-convolved pair integral (Section 3.7) and the mesh-only path is recovered by omitting the correction (the default).

Switching rule. The correction is applied per (source, node) pair, and a pair is *near* when the source

and node lie within R mesh nodes of each other *on every axis* (a per-axis box, not a Euclidean sphere; Fig. 1b). The output is tiled for parallelism and the near set of a tile is every source within R of the tile box on every axis, so the effective reach is R plus at most half a tile at a tile boundary. R defaults to 3 nodes and grows with the requested tolerance, $R = \mathrm{clamp}\big(3 + \lceil \log_{10}(10^{-3}/\mathrm{tol}) \rceil,\, 3,\, 6\big)$, and is user-overridable; it must cover the TSC support of 1.5 nodes.

Seam. The subtraction $F_{\mathrm{mesh}}(\mathrm{near})$ is a per-tile local zero-padded FFT (padding 2) of exactly those near sources (Section 3.7); because the convolution is linear and uses the same kernel as the global mesh, it equals the mesh value those pairs would otherwise receive, so the switch is exact at the seam.

Singular pairs. Where the observer falls inside a source's TSC support, the pair integral is split into an exact O(1) box term and a smooth remainder (Section 3.7). The box is the divergence-theorem field of the unit-density support box, an eight-term sum of the asinh primitive per component; the remainder is a tensor Gauss–Legendre integral of the subtracted integrand over the TSC knot panels (default order 5). For a moving source the box is γ-stretched along the motion for axial $\boldsymbol{\beta}$; an oblique $\boldsymbol{\beta}$ is integrated directly in the Lorentz-contracted frame by a ray integral aligned with $\boldsymbol{\beta}$.

Accuracy and cost. The mesh alone reaches 3×10^{-4} only at 8 cells ($\gamma = 1$) and 48 cells perpendicular ($\gamma = 10$) (Section 5.4), so the shell must extend past the 1.5-node support to where the mesh degrades; inside it the field is exact to roundoff and the evaluator's own error is 3.7×10^{-4} (stationary), 1.4×10^{-4} (axial), 1.9×10^{-3} (oblique, measured limit) (Section 5.6). The mesh costs $\mathrm{O}(M^3)$ in the padded grid and is independent of the particle count N; the correction costs one closed-form evaluation per near pair, of which a sparse bunch has $\sim N(2R)^3$, so R and the quadrature order are the cost levers ($R = 1.5 \to 4$ multiplies the wall by ≈ 3.9 at fixed N and grid, Section 5.7). A second switch occurs at the compute-box boundary: where the global output domain extends beyond the box, the margin is filled with an asymptotic moving multipole rather than the FFT (Section 3.8).

3.2 Deposit

Each source charge is deposited onto the zero-padded M^3 grid with the TSC shape, at the cell-unit spline argument $W((\mathbf{x}_i - \mathbf{x}_p)/\mathbf{cell})$, which is correct for any h. The result $\rho[i] = \sum_p q_p\, W(ih - \mathbf{p}_{\mathrm{cell}})$ lives in the N^3 sub-box of the padded grid, indexed with M-stride; the rest is zero. A fixed reduction order is required for MPI-layout invariance; the single-rank path uses `atomicAdd`.

3.3 Kernel sampling

The six TSC-convolved kernels are sampled on the same padded M^3 grid, one thread per FFT point, with the centered (negative-wrapped) index $d_a = a$ for $a < M/2$ else $a - M$. For each of the six components c, the kernel is evaluated at the physical separation

$$\mathbf{R} = (d_a\, h + \delta_{c,a})\, \mathbf{cell}, \tag{8}$$

where δ_c is the Yee staggered offset (Section 3.5). The kernel is therefore sampled at the same node spacing h as the deposit, with the shape one grid cell of its own axis wide. Non-finite samples are set to zero to avoid $0 \times \infty$ in the FFT.

3.4 FFT convolution and normalization

The convolution is a free-space, zero-padded FFT convolution, with the discrete convolution theorem in the unnormalized form of the vendor documentation [Press et al., 1986, NVIDIA Corporation, 2024]. The plan is built once, unnormalized, and reused for all thirteen transforms: one forward for ρ, six forwards for the kernels, and six inverse transforms. The unnormalized round trip is scaled by M^3; the read-out divides by $h^3 M^3$, because the sampled convolution evaluates the continuum integral of Eq. (1) by a Riemann sum whose nodes are spaced h cells apart, so each sample covers a physical volume h^3 (in cell units).

Padding. The grid is zero-padded to pad ≥ 2 so that the cyclic convolution does not wrap [Press et al., 1986]. With pad $= 1$ the cyclic convolution wraps and a 9.275×10^{-3} periodic-image artefact appears for a single stationary particle in a 20^3 box; at pad ≥ 2 the error is at machine precision ($\leq 10^{-15}$). The artefact grows for beamed sources: at pad $= 1$ the 100-source $\gamma = 10$ cloud shows a 7.5×10^{-1} worst-component error (Appendix B).

DC and Nyquist. For the R2C/C2R half-spectra the DC bin and the seven self-conjugate even Nyquist corners must be exactly real; the intermediate Nyquist planes are Hermitian-paired, not real. The R2C/C2R path matches the C2C path to 1.665×10^{-16} over all six components (Appendix B).

3.5 Read-out

The six field components live at the staggered Yee positions, in cell units:

$$\begin{array}{llllll} E_x & (0.5,0,0) & E_y & (0,0.5,0) & E_z & (0,0,0.5) \\ B_x & (0,0.5,0.5) & B_y & (0.5,0,0.5) & B_z & (0.5,0.5,0). \end{array} \tag{9}$$

The kernel is sampled at these offsets, so the convolution is already evaluated on the Yee grid. The read-out compacts the M^3 FFT result to the N^3 bunch box and applies the h^3/M^3 scale and the SI factor (Section 2). No further interpolation is needed.

3.6 Scale invariance

The kernel is homogeneous of degree -2, the deposit shape width equals the node spacing, and the Yee offsets are in cell units, so every relative metric depends only on r/h, the distance in cell units. At a fixed distance in cells the error is independent of h: refining h with the distance held in cells changes nothing, and the measured relative errors at $h = 1,\ 1/2,\ 1/4,\ 1/8$ are identical to machine precision for $\gamma \in \{1, 10, 100\}$, aspect 1 and 4. At a fixed physical distance r, refining h increases r/h and does reduce the error, at the cost of M^3 FFT work in $1/h$. The controlling length for a beamed source is the beaming cone width d/γ: where that width is sub-cell, the error is set by the cone, not by the mesh (Section 5.5).

3.7 The near shell

The default output is the whole mesh. An opt-in near shell replaces the mesh where the source and node are within R mesh nodes of each other on every axis (a per-axis box, not a Euclidean sphere; Section 3.1, Fig. 1b):

$$F = F_{\text{mesh}}(\text{all}) - F_{\text{mesh}}(\text{near}) + F_{\text{near}}(\text{near}), \tag{10}$$

where the subtraction uses a per-tile local zero-padded FFT of the near sources (padding 2) and the near term is evaluated per (source, node, component) by one of two backends. `exact` is an adaptive FP64 TSC-convolution quadrature shared with the independent reference driver; `fast` is the closed-form evaluator below. R is in mesh nodes and defaults to $R = \text{clamp}(3 + \lceil \log_{10}(10^{-3}/\text{tol}) \rceil,\ 3,\ 6)$ when the user leaves it unset; the default mode is `none`.

The fast evaluator splits the in-support integral of Eq. (1) into a singular part and a smooth remainder. The singular part is the eight-term asinh box of the kernel's divergence, O(1) per pair, built from the double integral

$$\begin{aligned} T(A,B,u) &= \int_0^A \int_0^B \frac{\mathrm{d}X\,\mathrm{d}Y}{\sqrt{u^2+X^2+Y^2}} \\ &= \begin{cases} A\,\operatorname{asinh}(B/A) + B\,\operatorname{asinh}(A/B), & |u| < 10^{-300}, \\ A\,\operatorname{asinh}\!\big(B/\sqrt{u^2+A^2}\big) + B\,\operatorname{asinh}\!\big(A/\sqrt{u^2+B^2}\big) & \\ \quad -u\,\arctan\!\big(AB/(u\sqrt{A^2+u^2+B^2})\big), & \text{otherwise.} \end{cases} \end{aligned} \tag{11}$$

with $A, B > 0$; each box face has 8 corners and 24 such calls per component, so the eight-term box is O(1) per pair. The form is even in u, so the $|u| \to 0$ branch must fire on $|u|$, not u (a negative u is a field point beyond a box face). The remainder is a plain tensor Gauss–Legendre quadrature on the subtracted integrand $(S - W^*)K$ over the TSC knot panels, each axis split at the TSC knots and at the anchor point s^*; the default order is 5 (8000 nodes per in-support pair). For a moving bunch the box is γ-stretched for the axial direction; the oblique direction instead uses a direct Lorentz-contracted (d-space) ray integral of $16 \times 32 \times 4$ nodes (angular rule polar-aligned with $\boldsymbol{\beta}$, four radial Gauss nodes per knot interval), with no remainder on that branch. Host and device use the same algebra; the full derivation is in Appendix A.

Assembly. Single-source tests at a tile centre and a tile boundary give an assembly residual of $2{\times}10^{-12}$ to $1.5{\times}10^{-11}$ SI ($4{\times}10^{-16}$ relative), flat with distance. A local FFT of the near set against the global FFT agrees at 10^{-15} for margins 0, 2, 4 and a 100-particle cloud. A direct reproduction with 100 particles at radius 3 gives an assembled field exact to 4.6×10^{-13} over the whole grid, while the plain mesh on the same plane is wrong by 5.6×10^{-1}. The fast backend changes no far number: the mesh-only field from the fast and the exact backends is bitwise identical, and the total field outside the near tiles is bitwise equal, so the tile subtraction stays intact. The seam jump excess at the fast/exact crossing, on a 16^3 grid, $R = 3$ cells, tile 4, cells 10^{-7} m, is $2.03{\times}10^{-10}$ per component for a stationary single particle ($2.8{\times}10^{-10}$ for the 64-source cloud).

3.8 The far-tail margin

For a global domain that extends beyond the compute box, the zero fill must not start until the field is below tolerance. The margin is placed at the smallest centre distance d where

$$E_{\text{tail}}(d) = \gamma_{\max}\, k \left(\frac{|Q|}{d^2} + \frac{2|\mathbf{p}|}{d^3} + \frac{3.5\,\|D\|_2}{d^4} \right) \le \text{tolerance} \cdot E_{\text{peak}}, \tag{12}$$

with Q the total charge, $\mathbf{p}$ the dipole, $D_{ij} = \sum q\,(3x_i x_j - |\mathbf{x}|^2\delta_{ij})$ the traceless quadrupole from the actual particle positions, $\|D\|_2$ its largest absolute eigenvalue, and $\gamma_{\max} = \max\sqrt{1+|\mathbf{u}|^2}$. Beyond a measured switch radius the writer evaluates an asymptotic moving multipole rather than a static one: the exact Liénard monopole of the total charge moving at the charge-weighted mean velocity, plus the rest-frame dipole and quadrupole evaluated at the Lorentz-stretched observation point and boosted back to the lab frame. For $\beta = 0$ this reduces exactly to the static multipole and for one particle it reproduces the point kernel in every direction; for an extended bunch it is the leading moving multipole, not a full Liénard multipole sum. The output file therefore has no interior zero fill. The margin is reported at run time as a bound in V/m; the switch error against an independent direct charge sum is in Appendix D.1, and the moving point-kernel check is `Margin.MovingFarTailMatchesPointKernel`.

Spatial far-tail multipoles. The tail can use one multipole per spatial supercell instead of a single global expansion. The bunch is split into boxes of N PIC cells (`-multipole-group-cells`, default 16; 0 keeps one global multipole), one set of moments is accumulated per occupied box about its own centre, and the tail sums the group expansions. For a two-cluster $\gamma = 10$ bunch the worst far-tail error against the direct Liénard sum at $6R$, $8R$ and $12R$ falls from 1.72×10^{-3} (global) to 1.81×10^{-4} (4-cell groups). The evaluation cost grows with the number of occupied groups per tail node, so the group size controls the accuracy-cost trade-off; large groups are cheaper, small groups are more accurate near the exact-box transition.

3.9 MPI parallelism

The program writes a global output domain larger than the compute box and computes the field in parallel across ranks. The written extent is set by `-global-grid` and `-global-origin`. The global array is never allocated: each rank owns a contiguous z-slab split into tiles of about 128 MiB, zero-filled, then filled with the compute-box overlap, the rest exact 0.0. All ranks open one openPMD series with `MPI_COMM_WORLD` and write disjoint `storeChunk` regions with collective flushes, so the output is a single dataset. A `-deterministic-deposit` flag gives a serial host deposit for bitwise comparisons.

The compute is parallel in two ways, neither of which is a domain-decomposed FFT. The mesh runs six independent component convolutions: every rank deposits and transforms the same ρ spectrum (replicated), rank r evaluates the components c with c mod size $= r$, and one `MPI_Allreduce` sums the six arrays. The near shell partitions the output tiles by linear tile index; each rank builds its own near-source list and local subtraction FFT for its tiles and writes only those corrections, and one `MPI_Allreduce` sums the correction arrays. The bunch is read once on rank 0 and broadcast, not redistributed. Because the mesh components are disjoint and each near correction element is owned by exactly one rank, the all-reduce order cannot change a value: the outputs are bitwise identical at 1, 2 and 4 ranks (Section 7.1).

4 Physical consistency

Five checks constrain the implementation against the physics it is built to reproduce. Each check names its measured value.

Liénard–Wiechert kernel. The kernel of Section 2 is the present-time Liénard–Wiechert field of a uniformly moving charge [Jackson, 2007]. The mesh reproduces the TSC-convolved Coulomb field of a single stationary particle to a worst relative error of 4.88×10^{-14}, E-component L2 residual 2.29×10^{-25}, and elementwise maximum 2.28×10^{-25}, at $d = 8, 16, 24, 32$ along all six axis directions, against the direct TSC stencil sum of the analytic Coulomb kernel (Coulomb anchor, Section 7).

Lorentz contraction. The field of a uniformly moving charge is compressed into a pancake of width d/γ perpendicular to the motion. The seven-velocity measurement (Appendix B, Fig. 10) tracks the pancake: at $\gamma = 10$ the mesh resolves the contracted field to 3.5×10^{-1} worst relative error on the mid-z plane of a 128^3 grid with 5×10^{-8} m cells, with the 5×5 source neighbourhood excluded (Appendix B), and the error grows with γ because the pancake becomes thinner than the mesh spacing (1.59 and 3.49 at $\gamma = 50$ and 100), not because of a solver error. Raising the resolution from 64^3 to 128^3 at the same physical domain leaves the relative errors unchanged.

Gauss's law. The discrete Gauss-law indicator, the Yee backward-difference div E residual normalised to the peak deposited ρ/ε_0, is 0.3534 for the mesh alone on the $N = 1000$ slice geometry ($\gamma = 1$) and 5.026 for $\gamma = 10$, against reference values 0.0904 and 0.0814. With the near shell enabled it moves to 0.0895 for $\gamma = 1$, reaching the reference value, and to 0.399 for $\gamma = 10$, a 12.6× drop that still leaves it 4.9× above the reference because the shell radius covers a fraction of the TSC support (Appendix C).

Yee staggering. The six field components are evaluated at the staggered positions of Section 3.5; the kernel sampling uses the same offsets, so the convolution is already on the Yee grid and the read-out applies no interpolation. The cross-validation against the PIConGPU `referenceOnly` field uses identical Yee offsets, grid, and origin, and the reference staggered point maps to a picBunch node by index with zero snap deviation (Section 6).

The $B = \boldsymbol{v} \times \boldsymbol{E}/c^2$ limit. The kernel satisfies $\mathbf{B}_n = \boldsymbol{\beta} \times \mathbf{E}_n$ exactly, and the mesh inherits it component by component. For a stationary bunch ($\gamma = 1$) every measured B component is exactly zero (Tables 8 and 17); for motion along $+y$ the component B_y is identically zero in the measured fields ($B_y \equiv 0$ in Figs. 13 and 10). Accordingly, the fully masked error panels in Figs. 6–12 (the whole B row at $\gamma = 1$; the B_y panel for $\gamma > 1$), Fig. 13 (B_y) and the B_x panel of Fig. 16 are fully masked, because the relative error of an identically zero reference component is undefined. The magnetic field at $\gamma = 10$ is about 3.3×10^{-9} of the E-dominated local field (Section 5), which is why the floored error metric is electric-only there (Section 5).

5 Validation and accuracy

Every error figure names its reference. The three references stay apart throughout: the mesh M, the TSC-convolved continuum T (Eq. (1), the quantity the mesh is built to compute), and the point field P (the un-smoothed Liénard–Wiechert field). M against T is the discrete deposit and convolution error. T against P is the TSC smoothing, an $O(1)$ effect at $\gamma h \gg 1$. M against P is the total.

5.1 The references and the metrics

The arbiter is an independent TSC-convolved continuum, a tensor Gauss–Legendre quadrature with a TSC-knot split and a singularity-cancelling power map; its quadrature error is below the 3×10^{-4} gate in every configuration measured here. Two more references constrain the method externally: an order-12 independent continuum (refine 3) against which the CLI is measured, and the PIConGPU `referenceOnly` field used in the cross-validation (Section 6).

Three metrics are reported and are not interchangeable; each claim in this paper names which one supports it.

Source-coincident nodes. The floored metric includes every node whose reference magnitude clears the floor, including a node that coincides with a source. The legacy `indfield.py` reference is unreliable at those nodes: its $p = 5$ power transform does not cancel the $1/r^2$ singularity, so it converges to a quadrature-dependent value (about 1.25% at the nearest source node). The hires validation of Appendix B therefore uses the campaign reference, which splits the singular part into an exact box field and

Table 1: Error budget. Each term is the measured contribution of one error source, with the measurement that establishes it. "Not measured" marks a term whose input data are not in the measurement campaign.

source	measured magnitude	statement
mesh / deposit	1.38×10^{-7} (M–T, $\gamma = 1$, aspect 1, $d = 48$); 2.6×10^{-4} ($\gamma = 10$, aspect 1, $d = 48$)	Table 2; far-field sweeps of Section 5.4
velocity binning	6.9×10^{-6} relative in the B_z peak ($B = 16$ vs 64, 5% energy spread); 3.9×10^{-4} relative worst pointwise in B_z for the 5% energy case; 1.4% worst against the arbiter for the 1 mrad case (not re-measured)	Appendix D, Table 17
near-source quadrature	3.71×10^{-4} stationary interior; 1.40×10^{-4} moving axial interior; 1.94×10^{-3} moving oblique interior (default observer-cell routing)	Table 6
far tail	8.53×10^{-2} V/m at the $6R$ switch, 1.87×10^{-3} of the tolerance level; 2.8×10^{-3} V/m worst on the 512^3 output (rosi_v100 direct-charge-sum run; comparison data not in the repository, not re-measured)	Appendix D.1, Table 18
output precision	float32 checkpoint: bitwise arrival under PICon-GPU, $\max_{\text{abs diff}} = 0.0$ after one step; float64 opt-in records a warning attribute	Section 6

integrates only the smooth remainder, so it is exact for the singular part. That table excludes the 5×5 nodes around the source and reports the full-plane value beside it, so the singular node is not hidden.

Floored. $|\text{ref} - \text{mesh}|/\max(|\text{ref}|, |\text{mesh}|)$, over points where $|\text{ref}| \geq 10^{-3}$ of the local (six-component) field peak. This is the relative per-component error of record for the electric field. For $\gamma = 10$ every B point is excluded by the floor: $|\mathbf{B}|$ is about 3.3×10^{-9} of the E-dominated local field, the value β/c at $\gamma = 10$ ($\beta = \sqrt{1 - 1/\gamma^2}$, $\boldsymbol{B} = (\boldsymbol{v}/c^2) \times \boldsymbol{E}$; the slice figures' colour bars agree, 2.3×10^4 T against 7×10^{12} V/m), far below the 10^{-3} threshold, so the floored metric reports no magnetic accuracy at all; the floored column of Table 3 is therefore electric-only at $\gamma = 10$.

Unfloored. The same ratio at every point. It is the only one that reaches the magnetic field, because the floor excludes it.

Absolute. $|\text{ref} - \text{mesh}|/\max_{\text{point}} |\text{ref}_c|$, the peak-normalized error; it supports whole-field statements.

The judge metric $|\text{ref} - \text{mesh}|/(1 + \max(|\text{ref}|, \text{mesh}|))$ is absolute-dominated: for fields of order 10^{-9} V/m the $+1$ denominator swamps the error, so it reports a crossover at $d = 2$ for every case and does not discriminate.

5.2 Error budget

Table 1 separates the measured error into its five sources, with the measured magnitude of each and the section that carries its full statement.

The mesh/deposit term dominates in the far field at moderate γ and is the only term that the mesh alone can reduce. The velocity-binning term is 6.9×10^{-6} at the B_z peak and 3.9×10^{-4} at the worst

Table 2: L2-pooled relative error at $d = 48$ cells, $h = 1$, perpendicular direction (L2 over the six components; single-component worst: $\gamma = 1$, aspect 4, $M-T$ 2.65×10^{-5}; $\gamma = 10$, aspect 1, 1.4×10^{-3}). $M-T$: mesh vs TSC-convolved continuum; $M-P$: mesh vs point field; $T-P$: TSC smoothing, mesh-independent. At $\gamma = 10$, aspect 1 the $M-P$ vs $M-T$ gap is the smoothing (L2 ratio 58×), not the mesh. Measured on `rosi_v100` (V100) and not re-measured in this campaign, while the $d \leq 32$ sweeps of Section 5.4 are on JURECA H100, so this $d = 48$ table and those sweeps are on different hosts).

γ	aspect	$M-T$	$M-P$	$T-P$
1	1	1.38×10^{-7}	9.89×10^{-8}	3.96×10^{-8}
1	4	2.65×10^{-5}	2.39×10^{-3}	2.40×10^{-3}
10	1	2.6×10^{-4}	1.5×10^{-2}	1.48×10^{-2}
10	4	3.33×10^{-2}	1.19×10^{-1}	1.44×10^{-1}
100	1	3.75×10^{-1}	2.20×10^{-1}	4.40×10^{-1}
100	4	2.78×10^{0}	2.49×10^{-1}	8.32×10^{-1}

pointwise node for the measured spreads at $B = 16$, and does not enter the single-velocity scope. The near-source quadrature term is the measured limit of the fast near backend in the oblique moving case (Section 5.6); the stationary and axial modes reach 10^{-3}. The far-tail term is bounded by the tolerance-level margin of Section 3.8 and is 1.9×10^{-3} of that level at the switch radius. Output precision is not a measured loss: the float32 checkpoint arrives bitwise in PIConGPU, and the float64 path is a uniform opt-in.

Shape bias. The deposited shape is a deliberate smoothing of the point field, not a mesh error. For a stationary point charge the relative bias at 2, 4 and 8 cells is -8.31×10^{-3}, -5.62×10^{-4}, -3.55×10^{-5} for CIC; -1.20×10^{-2}, -8.41×10^{-4}, -5.33×10^{-5} for TSC; and -1.55×10^{-2}, -1.12×10^{-3}, -7.10×10^{-5} for PQS. CIC is the closest of the three to the point field at these distances, TSC next, PQS the widest. The shape option exists to match the PIConGPU species, not to reduce this bias.

5.3 Reference-separated error at 48 cells

Table 2 shows two things. At $\gamma = 10$, aspect 1, the mesh against the point field ($M-P = 1.5 \times 10^{-2}$) is about 58× the mesh against the TSC-convolved continuum ($M-T = 2.6 \times 10^{-4}$, an L2 ratio); the difference is the TSC smoothing, which $M-P$ inherits and $M-T$ does not. A mesh-versus-point comparison measures the shape the deposit applies, not the discretization error. Second, the $M-T$ error is not converged perpendicular to the beam at $\gamma h \gg 1$: at aspect 1 it is acceptable only for $\gamma = 10$ and not for $\gamma = 100$; at aspect 4 it is not converged at any γ. Only the axial direction is clean at every γ.

5.4 Far field

Figure 2 plots the mesh against the TSC-convolved continuum. The plotted series is single-particle; the values are tabulated in Table 3. For the single-particle series the measured value is 1.1×10^{-4} at $d = 12$ for $\gamma = 1$, aspect 1, and 6.3×10^{-3} at $d = 32$ for $\gamma = 10$, aspect 1; at $\gamma = 100$, aspect 1 the relative error is 4.6×10^{-1} at $d = 32$. The aspect 4 series coincide with the aspect 1 series to machine precision at padding 1, where the error is set by the discretization and kernel rather than periodic wraparound at these distances. These are floored electric values; the magnetic field is excluded by the per-component floor at every swept γ. The window starts at $d = 2$, so the TSC-support interior ($|s| < 1.5$ cells) is not covered by these sweeps; it is what the near shell measures (Section 5.6).

Direction lattice. The fresh measurement lattice includes the axis, the face diagonals and the cube diagonals. At $d \geq 16$ the worst-case node can lie off-axis, so the tabulated floor is lower than the legacy axis-only value. The difference is the sampled lattice, not a mesh improvement.

Table 4 gives the first distance at which every larger distance is at or below the gate, per axis direction.

At $\gamma = 1$, aspect 1, the mesh is within 3×10^{-4} of both references from $d = 8$; the two are interchangeable there. From $\gamma = 10$ up, against the TSC-convolved continuum the perpendicular crossover is $d = 48$ ($\gamma = 10$, aspect 1, 2.6×10^{-4} L2, Table 2); the perpendicular crossover distance for the same configuration against the point field exceeds the measured window (> 48, Table 4). For the remaining

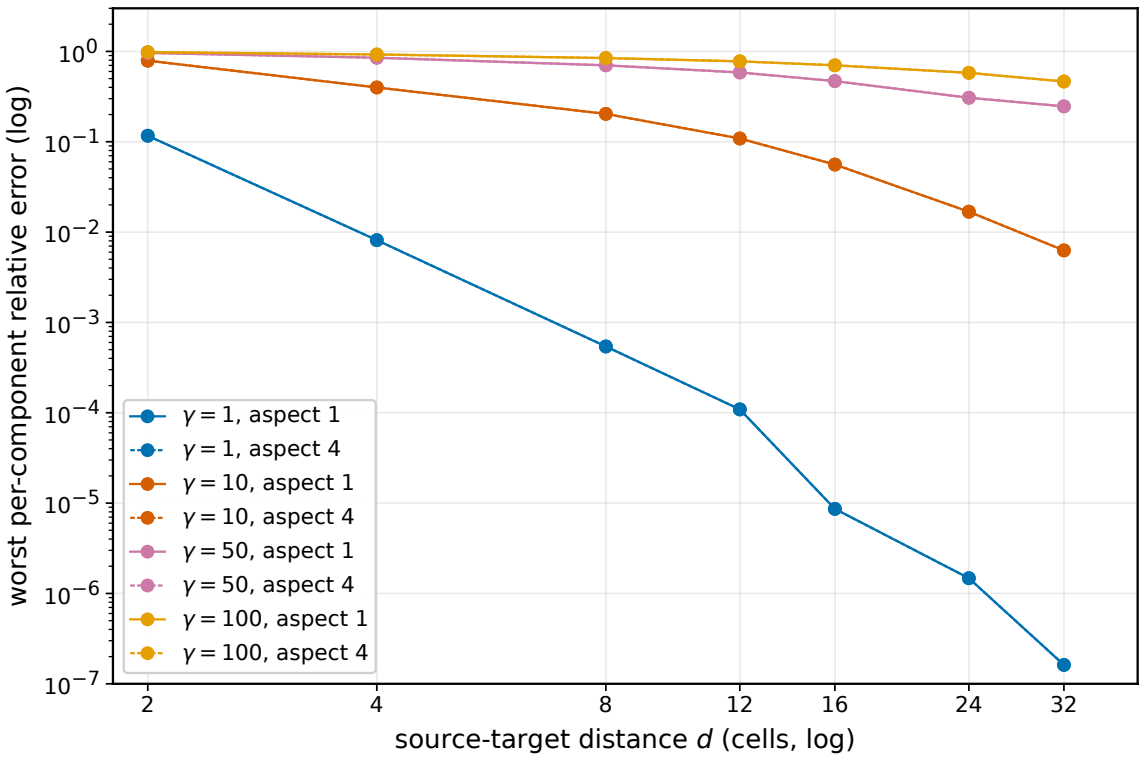


Figure 2: Worst per-component relative error against source-target distance d for the single-particle floored electric series, both axes logarithmic; hue is γ (blue $\gamma = 1$, vermillion $\gamma = 10$, purple $\gamma = 50$, orange $\gamma = 100$) and line style is aspect (solid 1, dashed 4), against the order-12 independent continuum reference (refine 3). At padding 1 the aspect 4 curves are numerically identical to the aspect 1 curves to machine precision, so the dashed lines coincide with the solid ones. The error falls from 1.2×10^{-1} at $d = 2$ to 1.6×10^{-7} at $d = 32$ for $\gamma = 1$, aspect 1. The single-particle values are in Table 3; the plotted window ends at $d = 32$. The $d = 48$ values are tabulated separately (Table 2) and carry a padding-1 wraparound floor, so they sit outside this window.

Table 3: Worst per-component relative error, per-component floor $|\text{ref}| \geq 10^{-3}$ of the component peak, over the electric components E_x, E_y, E_z, of the CLI against the TSC-convolved continuum, at $d = 16$ and $d = 32$ cells; the magnetic field is excluded by the floor. Order-12 independent continuum reference (refine 3), `jureca-h100`. Single-particle series.

γ, aspect	d=16	d=32
1, 1	8.6e−6	1.6e−7
1, 4	8.6e−6	1.6e−7
10, 1	5.6e−2	6.3e−3
10, 4	5.6e−2	6.3e−3
50, 1	4.7e−1	2.5e−1
50, 4	4.7e−1	2.5e−1
100, 1	7.0e−1	4.6e−1
100, 4	7.0e−1	4.6e−1

configurations the measured value at $d = 48$ exceeds the gate. The axial direction is 6 against the continuum at every γ.

5.5 The γ landscape: 10 to 100

The error landscape for a single particle at $\gamma = 10, 50, 100$ is measured on a 96-cell box, cell 10^{-7} m, aspect 1, axial, at $h = 1$ and $h = 1/4$ against the independent continuum. The metric is the peak-normalised absolute error per component at $d = 8$ cells for the h-comparison.

The error grows steeply with γ: 0.26 at $\gamma = 10$, 2.3 at $\gamma = 50$, 5.3 at $\gamma = 100$, at $h = 1$ and $d = 8$ cells.

Table 4: Crossover radius in cells: first d at which every larger d is at or below the gate, $h = 1$, L2 over the six components, per axis direction, worst over the three positive axis directions (>48: beyond the measured window). T: TSC-convolved continuum; P: point field. Oblique directions are $\sim 1.5\times$ larger. Measured on `rosi_v100` (V100) and not re-measured in this campaign; like Table 2 this is a V100 value, while the $d \leq 32$ sweeps of Section 5.4 are on JURECA H100, so the two are on different hosts.

	3×10^{-4}		10^{-3}			
γ	aspect	ref	axial z	perp x	axial z	perp x
1	1	T	8	8	6	6
	1	P	8	8	6	6
1	4	T	6	24	6	16
	4	P	48	> 48	32	> 48
10	1	T	6	48	6	48
	1	P	> 48	> 48	32	> 48
10	4	T	6	> 48	6	> 48
	4	P	> 48	> 48	32	> 48
100	1	T	6	> 48	6	> 48
	1	P	> 48	> 48	32	> 48
100	4	T	6	> 48	6	> 48
	4	P	> 48	> 48	32	> 48

Table 5: Single-particle mesh error at $d = 8$ cells, 96-cell box, cell 10^{-7} m, aspect 1, axial, `jureca-h100`. Peak-normalised absolute error, worst per component, at $h = 1$ and $h = 1/4$. The h-refinement exponent p (error $\sim h^p$) is a three-point fit over $h = 4, 1, 1/4$ cells on the $d = 2$ ring; it decreases from 1.33 at $\gamma = 10$ to 1.09 at $\gamma = 100$ ($R^2 = 0.979$–0.995).

γ	$h = 1$	$h = 1/4$	slope p	note
10	0.26	6.3×10^{-3}	1.33	
50	2.3	0.33	1.12	$h = 1$ off by more than the peak
100	5.3	0.90	1.09	$h = 1$ off by more than the peak

Table 6: Fast near evaluator, worst per-component relative error against the order-12 independent continuum reference, cells 10^{-7} m, CPU pointwise evaluator (host-independent; no CUDA/FFT linked). Regions: interior (inside the TSC support), face (on its boundary), outside (1.6–3.6 cells). The "remainder order" column is the Gauss–Legendre order of the smooth remainder; the singular asinh box is O(1) in every row (moving axial uses the γ-stretched box; moving oblique uses the direct d-space ray integral, $16 \times 32 \times 4$ nodes, no remainder).

γ, direction	remainder order	interior	face	outside
1, stationary	5	3.71×10^{-4}	1.96×10^{-5}	1.63×10^{-7}
10, axial	8	1.40×10^{-4}	1.11×10^{-4}	2.12×10^{-6}
10, oblique	5	1.94×10^{-3}	3.94×10^{-3}	1.01×10^{-4}

Refining to $h = 1/4$ brings these to 6.3×10^{-3}, 0.33, and 0.90. The controlling length is the beaming cone width d/γ: for $\gamma \geq 50$ the cone is sub-cell, so sub-cell resolution of the cone is the binding limit. In the node-unit kernel the h-refinement exponent p (three-point fit over $h = 4, 1, 1/4$ cells on the $d = 2$ band) is 1.33, 1.12, and 1.09 at $\gamma = 10, 50, 100$ ($R^2 = 0.979$, 0.990, 0.995); it decreases mildly with γ, so refinement reduces the error across the range while the fitted order does not rise. At $d = 8$, $\gamma = 100$ the cone is $d/\gamma = 0.08$ cells and a further $h \leq$ that cone width is what 10^{-2} requires. These are measured limits of the mesh at a fixed distance in cells.

5.6 The fast near evaluator, per component, by region

Table 6 gives the domain of validity of each near-shell mode. The metric is $\max|\text{mesh}-\text{ref}|/(1+\max|\text{ref}|)$ over 20 interior, 16 face and 16 outside points, cells 10^{-7} m. The per-component values behind the stationary interior row are 3.71×10^{-4}, 1.80×10^{-4}, 9.69×10^{-5}, 0, 0, 0 for $E_x, E_y, E_z, B_x, B_y, B_z$ (the magnetic field is identically zero for a stationary bunch); the moving axial interior row is 1.40×10^{-4},

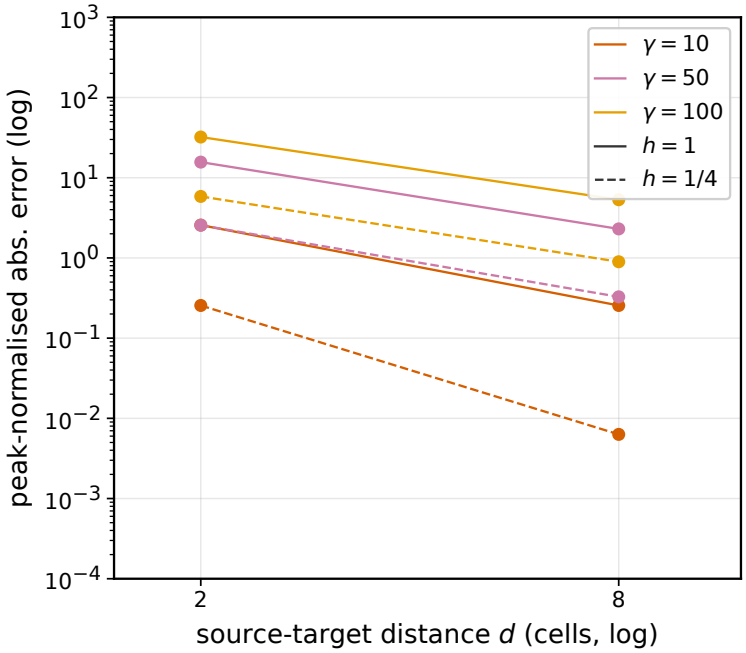


Figure 3: Peak-normalised absolute error against distance on log-log axes, single particle, 96-cell box, cell 10^{-7} m, aspect 1, axial, `jureca-h100`. Solid/dashed is $h = 1/1/4$ and hue is $\gamma = 10$ (vermillion), 50 (purple), 100 (orange). The plotted distance range is 2 to 8 cells; larger distances are limited by the finite box size at the tested resolution, so they are not shown. At $h = 1$ the error exceeds the peak field by $d = 8$ cells for $\gamma = 50$ and $\gamma = 100$.

1.35×10^{-4}, 2.20×10^{-5}, 1.35×10^{-4}, 1.40×10^{-4}, 0; the moving oblique interior row is 7.01×10^{-4}, 1.94×10^{-3}, 9.78×10^{-4}, 6.44×10^{-4}, 4.15×10^{-4}, 5.54×10^{-4}. The interior entries are below the published values, because the evaluator (`obaRayField` oblique ray integral) and the reference (a refined `indfield`) of this table are later revisions of the published ones; the difference is largest for the moving oblique interior, 8.15×10^{-2} in the published table, 1.94×10^{-3} in this table. The shipped default keeps the base order and ray and applies the exact evaluator only to the observer-containing TSC support cell; it lowers the stationary interior row from 6.82×10^{-4} to 3.71×10^{-4} and leaves the moving axial and oblique rows unchanged. A tolerance of 10^{-4} or below, or a tile with $\gamma \geq 100$, selects the full-support exact path.

The validity of each mode follows from the table. `exact` is the adaptive FP64 quadrature shared with the reference driver. Its in-support moving branch integrates in Lorentz-contracted coordinates, where the Liénard kernel becomes a Coulomb kernel and the beaming anisotropy moves into a smooth TSC density; the pyramid singularity cancellation is then exact for any velocity. Against the order-12/refine-3 arbiter at $\gamma = 10$ it reaches about 10^{-4} at the axial points $(1, 0, 0)$ and $(0, 0.5, 0)$, about 10^{-3} for the oblique points, and about 3% at the closest axial node. It is therefore a usable reference, but the closest node is not a blanket 10^{-3} path. `fast` meets 10^{-3} in every region for the stationary and moving axial modes; its measured limit is the moving oblique interior, 1.94×10^{-3} at the default order. Against the order-12/refine-3 continuum arbiter the fast backend is 1.6% at the closest moving in-support node. The driver panels in Appendix C are an upper bound on the in-support error (Section C).

The shell radius and order are the cost levers (Section 5.7). The fast evaluator runs the same algebra on the CUDA device when the build has one and a device is present (`-near-shell fast`); `fast-gpu` is the explicit device request. The kernel uses one thread per (pair, component) and agrees with the host path to 5.4×10^{-16} to 8.0×10^{-16}. The near-shell radii used on the slice geometry are stated in mesh nodes of that $h = 1$ grid (Appendix C, shell-radius unit note). On `rosi_v100` (64^3 output, $R = 3$, tile 4, 2000 particles, 6.144×10^6 pairs) the final default runs in 0.710 s on the device against 12.77 s on the host, a 18.0× speedup; the pre-fix base is 0.704 s against 13.26 s. Both are one idle V100 node, so the ratio is a same-node comparison. It falls back to the host with a warning when no device is present.

Table 7: Measured CLI wall at 2×10^6 particles, 20-cell bunch, $66 \times 68 \times 66$, cell 2.5×10^{-7} m, tile 4, one H100 (JURECA H100, `sm_90`, ADIOS2-BP5, CUDA 13.4). "near-shell wall" excludes the mesh, "end-to-end wall" includes it. The mesh-only wall is the mesh stage of the $R = 1.5$ row: 0.500 s of its 500.95 s end-to-end wall, of which the compute is 0.083 s.

configuration	near-shell wall (s)	end-to-end wall (s)	worst near error
mesh only (mesh stage of the $R = 1.5$ row)	0	0.500	no near correction
support-only shell, $R = 1.5$, order 5	498.24	500.95	7.99×10^{-4}
three-cell shell, $R = 3$, order 5	1265.62	1268.29	7.99×10^{-4} in, $\leq 1.6 \times 10^{-7}$ out
four-cell shell, $R = 4$, order 5	1956.61	1959.31	7.99×10^{-4} in, $\leq 1.6 \times 10^{-7}$ out

5.7 Cost at 2×10^6 particles, one H100

The[1] near shell is the dominant term in every shell row. The mesh-only wall is the mesh stage of the $R = 1.5$ row: 0.500 s of its 500.95 s end-to-end wall, of which the compute is 0.083 s (deposit 37.2 ms, kernel 1.1 ms, FFT 13.2 ms, read-out 18.2 ms, writer 13.1 ms; the per-repeat values are in `docs/measurements/cost_meshonly_bp5_20261108.csv`); the rest of the row is the near shell. A standalone mesh-only run of the same geometry was not re-measured in this campaign, so the mesh-only row is reported as a component of the $R = 1.5$ run, not as a separate measurement. The support-only shell costs 500.95 s end-to-end at 7.99×10^{-4} interior error, the three-cell shell 1268.29 s, and the four-cell shell 1959.31 s, at the same interior error. $R = 1.5$ to $R = 4$ multiplies the end-to-end wall by 3.9× at fixed particle count and grid, while the accuracy of the in-shell region is already at 7.99×10^{-4} at $R = 1.5$. The walls predate the hybrid near path; the final default routes only the observer-containing support cell to the exact evaluator, and on the 64^3, $R = 3$ workload that policy is 0.91–0.96× the base CPU wall, 1.006–1.009× the base on-node GPU wall, and 1.16–1.30× the base GPU wall for the 12-cell cloud. A stationary integrated rerun was not made, so Table 7 keeps the measured pre-hybrid walls and its error column. The shell rate on the integrated path is 0.87×10^6, 1.15×10^6 and 1.36×10^6 pairs/s for the three radii, about 2× slower per component than the standalone kernel, the overhead of the per-tile local FFT and launch in the integrated path. The measured per-pair penalty is 1.79× at $R = 1.5$, 1.40× at $R = 3$ and 1.29× at $R = 4$ against the published rates, so the gap is concentrated at small R rather than uniform. At $R = 1.5$ it splits into two parts on this host and stack: the v0.9.2 binary is 1.42× slower than the v0.9.1 binary (near-shell wall 498.24 vs 352.30 s) on the same node, input and pair count. The difference between the two versions is the oblique near-shell convergence fix: its ray integral raises the near-shell kernel arithmetic 1.44× at unchanged occupancy, while the co-changed stencil term adds $< 1\%$, and it is what improves the moving-oblique interior accuracy from 8.15×10^{-2} to 1.94×10^{-3} (§5.6). The residual host/stack term: the v0.9.1 binary rebuilt on this stack is 1.26× the published value (end-to-end 353.47 vs 281.0 s), the residual host/stack difference. A clock probe over the whole block (idle 345 MHz, load 1755 MHz, the maximum SM clock, 200–234 W, 60–63 C, throttle reason $0x0$) rules out throttling; neither term is a clock effect. A distributed FFT decomposes the output field in z slabs, one slab per rank. The end-to-end field, with the device TSC deposit into the interleaved complex slab, matches the replicated field to 3.2×10^{-16} at one rank and 5.3×10^{-16} at four ranks in the written file on HIP+MPI (hal); the same four-rank case on CUDA+MPI (jureca-h100, four H100s) gives 2.5×10^{-16}. The host FFTW backend supports strided plans in the current alpakaVendor and matches the replicated field to 1.9×10^{-16} at two MPI ranks. The automatic policy selects the slab path for every multi-rank run (single-rank auto keeps the replicated path below its padded-node threshold). At a 640^3 output the per-rank host memory is 41.1 GB at 2 ranks and 20.6 GB at 4, the quantity the decomposition buys: the output field is slab-local, while the bunch remains replicated on every rank. The `MPI_Alltoallv` transposes are 74–85% of the FFT-plus-transpose time, so the distributed path is a memory feature rather than a speed feature.

5.8 Determinism

Repeated runs of the CLI are bit-identical: two runs of the same input produce byte-identical E/B mesh payloads (md5 `d3641d64...`, 15 092 736 bytes). The CLI also rejects a below-floor `-tolerance` (10^{-15})

[1] The source drawer is reconstructed, not recorded: `gen_inputs.py` regenerates the 20-cell bunch at 2.5×10^{-7} m cell spacing centred in the $66 \times 68 \times 66$ grid, which fixes the grid origin at $(-5.75, -6.0, -5.75) \times 10^{-6}$ m; the original origin was not recorded. All field-solver settings (cell size, grid shape, tile, order, near-shell radius, backend) are the recorded campaign ones.

Table 8: Whole-field comparison of picBunch against the PIConGPU `referenceOnly` field, $N = 1000$, $32 \times 32 \times 16$, cell 2×10^{-7} m, 16384 nodes/component, rosi_v100. `worst_rel` is peak-normalized; median and p99 are per-point over nodes with $|\text{ref}| > 10^{-3} \times$peak; L2 is pooled. Not a truth comparison.

component	worst_rel	median	p99	L2
$\gamma = 1$ (B identically 0)				
E/x	1.48×10^{-1}	3.92×10^{-3}	3.20×10^{-1}	4.62×10^{-2}
E/y	1.49×10^{-1}	3.76×10^{-3}	2.93×10^{-1}	4.75×10^{-2}
E/z	1.22×10^{-1}	7.81×10^{-3}	3.65×10^{-1}	5.67×10^{-2}
$\gamma = 10$ (B/x identically 0)				
E/x	1.36×10^{-1}	4.80×10^{-3}	1.23×10^{-1}	4.44×10^{-2}
E/y	1.98	1.36×10^{-2}	7.25	7.63×10^{-1}
E/z	1.93	3.33×10^{-2}	10.1	9.19×10^{-1}
B/y	4.42×10^{-1}	3.12×10^{-2}	1.41	2.58×10^{-1}
B/z	4.63×10^{-1}	1.52×10^{-2}	0.93	2.20×10^{-1}

with a nonzero exit. The `gridUnitSI` is applied once, in the reader, and the mesh geometry follows `gridSpacing` · `gridUnitSI`; a mesh without a positive `gridUnitSI` is rejected.

5.9 Anisotropic cells

At anisotropic spacing the accuracy is quantified per geometry, not inherited from the cubic case. On a grid with per-axis cell spacing $(1, 1, 4)$ the stretched z axis matches the continuum to $\approx 10^{-3}$–10^{-7} across the in-grid distances, the same scale as the cubic control; the short x/y axes sit in the near-field TSC regime at a short physical distance ($d = 4$ cells is only 4 m along x/y, so $E_z \approx 0.19$ stationary and E_x/E_y up to ≈ 0.9 at $d = 4$ for the $\gamma = 10$ case). The mesh is not wrong because the spacings differ; the near points on the short axes are physically closer to the source.

6 Cross-validation and checkpoint transport

Two independent checks constrain the method externally: a cross-validation against PIConGPU's exact-quadrature field reference, and a checkpoint and restart transport. Neither is a truth comparison.

6.1 Cross-validation against the PIConGPU reference

The reference is the TSC-convolved field (Section 5), evaluated in PIConGPU [Burau et al., 2010] in its reference-only field-initializer mode. The comparison passes for every sampled point, and both sides use identical Yee offsets, grid, and origin; the reference staggered point maps to a picBunch node by index with zero snap deviation.

At 318 staggered points on a $64 \times 64 \times 32$ grid, cell 10^{-7} m, the overall worst per-component relative error is 4.86×10^{-2} for $\gamma = 1$ (E/z) and 6.20×10^{-1} for $\gamma = 10$ (E/z; E/y is 5.16×10^{-1}), against a reference reproducibility floor of order 10^{-6} (reported from the `rosi_v100` campaign, original provenance the PIConGPU `referenceOnly` run; the comparison file is archived outside this repository and was not re-measured in this campaign). picBunch is about four to six orders of magnitude above that floor; the measured error is concentrated at the core and the beamed peaks, where the TSC support is, while the bulk agrees to below 1–3% and B is exactly zero for $\gamma = 1$ in both components.

The whole-field comparison, $N = 1000$ on $32 \times 32 \times 16$, cell 2×10^{-7} m, 16384 nodes per component, gives the same picture: $\gamma = 1$ bulk median 3.8–7.8×10^{-3} with an L2 of 4.6–5.7×10^{-2}; $\gamma = 10$ bulk median 1.4–3.3×10^{-2} but a beamed-core overshoot that pushes the E/y, E/z L2 to 0.76–0.92. The forward E/x is well resolved at $\gamma = 10$ ($\sim 6\%$); the beamed transverse E/y, E/z are 50–130% worst at the core, where the Lorentz-contracted field is a sub-cell pancake below the TSC support and the spectral near-field kernel resolution.

The disagreement is per-source, local to the TSC support, and independent of particle count: on a $64 \times 64 \times 32$, cell 10^{-7} m, stationary bunch ($\gamma = 1$) the band is flat in N from 1 to 1000 (10^{-3}–3×10^{-2}). The $O(1)$ values at $N \geq 10$ occur only because the global worst sample point falls 0.415–0.65 cells from a source, inside that source's TSC support; an incorrect charge weighting would grow with N, and it does not.

6.2 Checkpoint format and bitwise arrival

The checkpoint writer emits a PIConGPU openPMD field checkpoint matching PIConGPU's own plugin schema: root layout, openPMD standard version 1.1.0 [Huebl et al., 2018] (written by the openPMD-api 0.17.1 software), E, B $x/y/z$ at the exact Yee positions, grid spacing and offset, and the `picongpu_idProvider` mesh metadata. The default precision is uniform `float32` for every field and particle record, so a stock single-precision PIConGPU can load the file; `float64` is a uniform opt-in that records a warning attribute because stock PIConGPU cannot load double records. Every species from the input bunch is copied, the `charge` and `mass` records are written rather than omitted, and the species is a single particle patch, `patch` = `"bbox"` or `"domain"`. For a multi-rank restart the writer splits the species into one patch per output rank: `patch_count` in `[placement]` gives cumulative, disjoint `numParticlesOffset` ranges, with one-cell overlap between neighbouring z regions, and every particle belongs to exactly one patch.

Field arrival is bitwise. A fresh checkpoint written from a two-species input (16 electrons and 16 ions, `picongpu_idProvider` startId 0, nextId 32) restarts in PIConGPU with exit 0 under both `patch` modes (`bbox`: offset $(7, 8, 3)$, extent $(8, 10, 6)$; `domain`: offset $(0, 0, 0)$, extent $(24, 24, 12)$), all six field components bitwise-equal, $\max_{\text{abs diff}} = 0.0$, electrons with ids 0 to 15, first-step displacement exactly $\beta c\, dt$ with $\beta = 0.9949874367$. The multi-rank restart is verified on the $24 \times 24 \times 96$ grid with the 4 ranks in z ($-d\,1\,1\,4$, per-rank $24 \times 24 \times 24$), which is an integer multiple of the per-rank grid on every axis and so needs no automatic grid adjustment: a 256-particle bunch written with `patch_count` 4 (four particle patches, one per rank, with two empty patches for this compact bunch) restarts PIConGPU with exit 0, and after one step all six field components are bitwise-equal to the checkpoint, $\max_{\text{abs diff}} = 0.0$ [picBunch project, 2026, measured in].

Placement and particle-data integrity. A 12-marker bunch with an asymmetric position and momentum pattern was placed at the xz middle and the y first quarter of a $32 \times 64 \times 32$ domain. picBunch reports centroid fractions $(0.500000, 0.250000, 0.500000)$; after one PIConGPU step the centroid is $(0.500755, 0.250446, 0.501466)$, the raw momentum records are preserved exactly, and the per-particle displacement matches $dt\, c\, u/\gamma$ to 4.1×10^{-7} cells. The displacement-to-momentum ratio is one positive scalar on all three axes (spread 8.9×10^{-6}), so the restart has no axis permutation and no mirror. One metadata requirement is sharp: PIConGPU restart ignores the particle `unitSI` and interprets momentum in `sim.unit.mass()`·`unit_speed`, so a checkpoint written with the electron mass as `unit_mass` restarts with all momenta 10^4 too large; `unit_mass` must equal the deck's `sim.unit.mass()`.

On a $24 \times 24 \times 12$ deck the restart is exit 0 with `-periodic 1 1 1`, and the 16 electrons load; initialization drops from the ~ 116 s neuralBunch path to 1.8 s because the checkpoint carries the field. For an open-boundary deck the writer copies the `Convolutional PML` slabs from the input checkpoint verbatim (converting a non-float32 slab to float32), so the PML state arrives with the restart. On a real LWFA checkpoint (638976 electrons, eight `Convolutional PML` groups with six datasets each, $64 \times 128 \times 64$, `-periodic 0 0 1`) the writer copies all eight groups and the restart exits 0; the copied PML datasets are float32 and bitwise equal to the source, initialization is 3.986 s, and the main E and B records are recomputed from the particles and stay float32.

7 Performance and scaling

All compute timings in Table 9 are on JURECA H100 (jureca-h100), four H100 PCIe 80 GB (`sm_90`), exclusive allocation, one GPU job at a time, GPU 0. Build: CUDA 13.4, ADIOS2 2.12.1 + openPMD 0.17.1, padding 2, double. Each configuration ran nine timed repeats after two untimed warm-ups; the SM clock was 1755 MHz (the GPU maximum) under load in every row. Stage values are medians; the total is reported as its median and, for the two bimodal rows, as two mode totals. The four independent $\gamma = 10$, $N = 2 \times 10^6$ repeats give total medians 51.88, 52.01, 51.69, and 52.63 ms, a spread of 1.8% of the minimum; the $\gamma = 100$ row is 51.93 ms, inside that spread. Each total is the sum of deposit, kernel, FFT, and read-out; device-buffer allocation and the cuFFT plan are excluded and reported separately (the plan is `plan_ms`; the allocation is `device_alloc_MiB`). The acceptance gate is IQR/median $\leq 2\times$ the measured B0 floor, which is 12.5% here; twelve of the fourteen rows pass it and are quotable, and the two that do not are bimodal and are reported as two modes (below). These compute-only totals exclude the checkpoint writer; the end-to-end totals with the writer are in Table 10, a different measurement path.

Table 9: Per-stage medians (ms) and device allocation (MiB), padding 2, double, JURECA H100 (jureca-h100) GPU 0, CUDA 13.4. "box" is the source-box size in cells; M is the padded FFT array per axis. Stage columns and the total are medians over nine timed repeats (two untimed warm-ups); the plan (`plan_ms`) is reported separately in `matrix.csv`. The two rows marked bimodal are not quotable as a single number under the protocol; their total is given as two mode totals (low/high) and the FFT/read-out that split are flagged with (b). The low mode is used for derived ratios (the typical case); the high mode is reported alongside it.

N	box	γ	h	n	M	total	deposit	kernel	FFT	read-out	dev
10^3	20	10	1	30	60	19.08	0.78	0.23	6.31	11.86	8.4
10^4	20	10	1	30	60	18.88	1.17	0.21	6.07	11.60	8.7
10^5	20	10	1	30	60	20.80	2.49	0.22	6.24	11.97	11.4
10^6	20	10	1	30	60	36.65	17.53	0.22	6.32	11.84	38.9
$2{\cdot}10^6$	20	10	1	30	60	51.88	33.42	0.20	6.13	11.79	69.4
$2{\cdot}10^6$	50	10	1	60	120	85.73	34.29	0.83	29.77	20.39	127.4
$2{\cdot}10^6$	100	10	1	110	220	115.99/241.89†	37.73	5.35	41.01/106.62(b)	21.69/81.98(b)	468.7
$2{\cdot}10^6$	20	10	1/2	50	100	78.69	33.71	0.53	24.13	20.07	99.5
$2{\cdot}10^6$	20	10	1/4	90	180	124.61/162.10†	36.29	2.97	66.66	11.93/49.42(b)	284.5

† Bimodal row (FFT and/or read-out split into two modes); the total is the sum of the stage cluster medians per mode, reported as low/high. The mode split is per repeat: the FFT and read-out of a given run land in the same mode.

Reproducibility. The matrix of Table 9 is the `picBunchBench` sweep, run with picBunchBench--sweep--repeat9--warmup2--csvmatrix.csv (two untimed warm-ups plus nine timed repeats, one process at a time). The per-row statistics, the B0 gate, the two bimodal rows and their mode splits, and the clock/sampler record are in `docs/measurements/matrix_h100_20261108.md`. The input bunches are not checked into the repository; the sweep configuration is the generator. The fast GPU kernel tools/closed-form/gpu_extract.cu is part of the source. The H100 runs used CUDA 13.4.

Correctness anchor. Every timing run is anchored by the Coulomb limit ($\mathbf{u} = 0$): one particle at a node, far nodes at $d = 8, 16, 24, 32$ along all six axis directions. The mesh field against the direct TSC stencil sum of the analytic Coulomb kernel has worst relative error 4.88×10^{-14}, E-component L2 residual 2.29×10^{-25}, and elementwise maximum 2.28×10^{-25}, passing the 10^{-12} threshold. This confirms that the zero-padded FFT convolution reproduces the TSC-convolved Coulomb field to roundoff (Section 4).

Figure 4 sweeps one axis at a time: the top panel fixes the 20-cell box and varies N, the bottom panel fixes $N = 2{\cdot}10^6$ and varies the box. The scaling is in two regimes, with the crossover set by the box, not the particle count. At large N the deposit dominates (0.78 to 33.4 ms on the 10^3 to $2{\cdot}10^6$ rows): the $10^3 \to 10^6$ leg is 22.5× (0.78 → 17.53 ms) and the $10^6 \to 2{\cdot}10^6$ leg 1.9× (17.53 → 33.42 ms), so it is sub-linear below $\sim 10^6$ (1.17, 2.49 ms at 10^4, 10^5) and approaches linear in N at the top of the sweep, where the FFT, kernel, and read-out stay small at the 20-cell box (FFT 6.1–6.3 ms, read-out 11.6–12.0 ms, kernel 0.20 ms). The total is flat across $\gamma = 1, 10, 100$ (51.52/51.69/51.93 ms, within 0.8% of the minimum). At large boxes the FFT and read-out scale with M^3: from $M = 60$ to $M = 220$ the total rises from 51.88 ms (20-cell) to 85.73 ms (50-cell, 1.65×) and to the 100-cell row, which is bimodal with low and high modes of 115.99 ms (2.24× the 20-cell total) and 241.89 ms (4.66×); the low-mode ratio is quoted for the typical case and the high mode is reported alongside it. Halving h costs 1.52× from $h = 1$ to $h = 1/2$ (51.88 → 78.69 ms) and, for $h = 1/4$ (bimodal, read-out), 1.58× in the low mode (124.61 ms, 2.40× the $h = 1$ total) and 2.06× in the high mode (162.10 ms, 3.12× the $h = 1$ total) because the deposit stays roughly constant (33.4–36.3 ms). The per-h totals are 51.88, 78.69, and 124.61/162.10 ms (low/high mode) for $h = 1, 1/2, 1/4$ (padded array $M = 60/100/180$, 2×10^6 particles, 20-cell box).

No configuration came near the 80 GB device limit (the H100 PCIe); the largest exact device allocation was 468.7 MiB. Memory scales with the padded bunch box, not the global field.

Table 10 is the same pipeline with the writer in the loop, on the V100 `rosi_v100`. The deposit grows with N (2.70 → 31.56 ms from 10^4 to 10^6), and the FFT is still box-bound and flat. The writer is the dominant added term: 11.29 ms at 10^4 and 7.54 ms at 10^6. The $N = 10^6$ total is 47.15 ms with the writer, versus 61.92 ms in the V100 compute-only sweep (the V100 row of Table 9, which is otherwise H100); the two campaigns differ in measurement path and timing boundary, so neither is the other's total plus the writer. The compute-only H100 total for $N = 10^6$ is 36.65 ms.

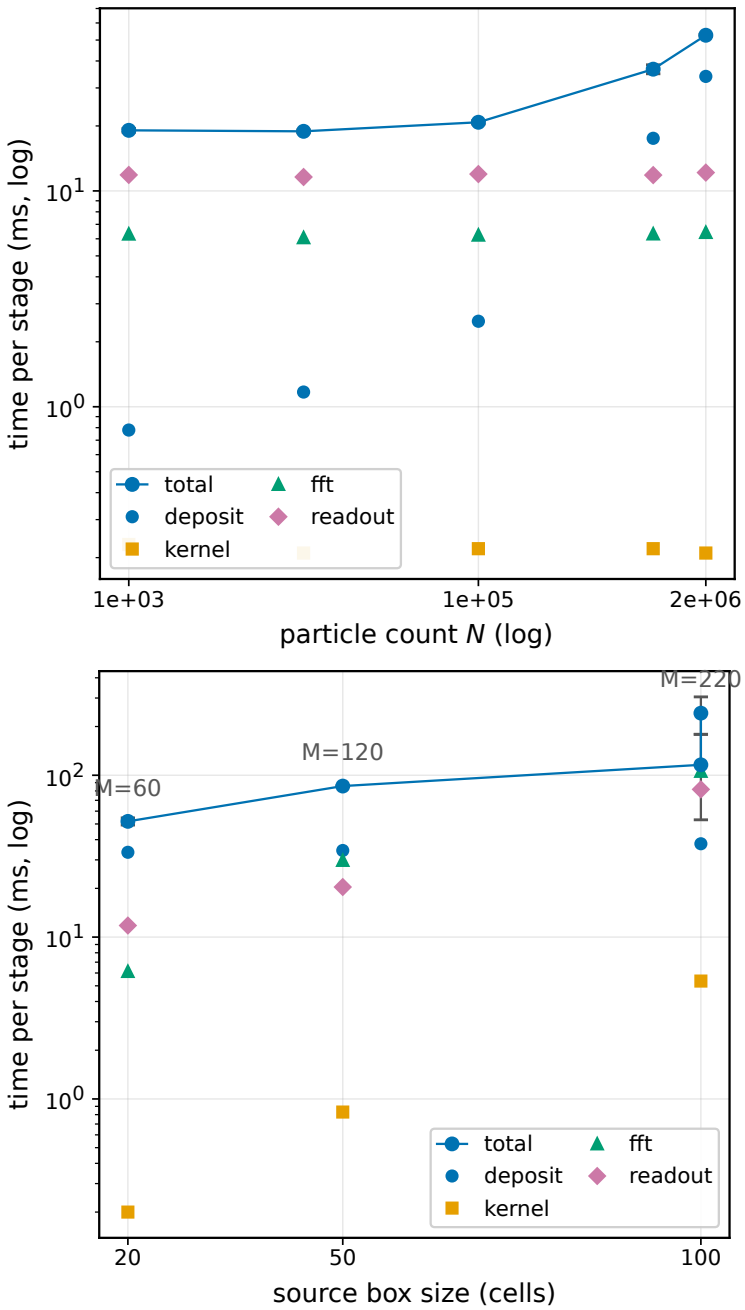


Figure 4: Pipeline time versus particle count (top, log-log) and source-box size (bottom, semi-log), $\gamma = 10$, $h = 1$, double precision, JURECA H100, CUDA 13.4. Blue is the total: the median over the kept repeats of Table 9 with an IQR bar; at $N = 2 \times 10^6$, box 20 the four independent repeats are all shown. The two bimodal rows plot both mode totals (low/high) instead of a single median. Markers are the per-stage medians. The deposit dominates at $N = 2 \times 10^6$ (33.4 of 51.88 ms), and the total rises with M as the FFT and read-out grow (51.88 ms at $M = 60$, 85.73 ms at $M = 120$, 115.99/241.89 ms at $M = 220$).

7.1 MPI rank scaling

With MPI the production policy is the slab-local distributed transform: `auto` selects it for every multi-rank run, and each rank keeps only its output z-slab. The measurements below force the same path at one rank so all three rows are comparable. At $N = 2 \times 10^6$, 256^3, with the ADIOS2/BP5 checkpoint, the total is 18.19, 10.08 and 7.23 s at 1, 2 and 4 ranks; the FFT is 2.31, 1.15 and 0.62 s and the distributed BP5 writer 0.71, 0.65 and 0.55 s. Fig. 5 plots these three rows. The distributed FFT scales to 3.73× at four ranks and the total to 2.52×; the writer scales sub-linearly, to 1.30× at four ranks, the remaining

Table 10: End-to-end pipeline time including the checkpoint writer and the openPMD reader, medians over 3 repeats, `rosi_v100` GPU 0 (V100), 20-cell box, $\gamma = 10$, CLI `-tolerance` 10^{-3}. rosi_v100 run: the raw per-repeat CSV is not in this repository and the rows are not in the current campaigns (the compute-only Table 9 is the H100 JURECA campaign). A different measurement path from the compute-only Table 9: the totals include the writer (the dominant added term); the per-stage split is deposit/kernel/FFT/readout/writer.

N	total	deposit	kernel	FFT	readout	write
10^4	22.00	2.70	0.53	4.53	3.06	11.29
10^5	24.87	5.82	0.56	4.31	3.80	10.65
10^6	47.15	31.56	0.46	4.20	3.95	7.54

Table 11: Distributed-mode strong scaling on JURECA H100 (4× H100 PCIe 80 GB) with the ADIOS2/BP5 checkpoint, medians of three serialized runs (one warm-up discarded). $N = 2 \times 10^6$, 256^3, one bin, tolerance 10^{-3}. The distributed path is forced at one rank and is the automatic path for the multi-rank rows.

	1 rank	2 ranks	4 ranks
total (s)	18.19	10.08	7.23
FFT (s)	2.31	1.15	0.62
writer (s)	0.71	0.65	0.55

HDF5 comparison, not plotted: the 2026-11-07 campaign (collective openPMD/HDF5 writer, one rank on the same node) gave totals 21.6/12.3/13.0 s with a collective writer of 4.19/2.88/6.46 s: the writer peaked at two ranks and regressed at four under collective-HDF5 contention on the shared series. The raw per-rank CSV of this campaign is not in this repository. The BP5 rows above are the plotted result.

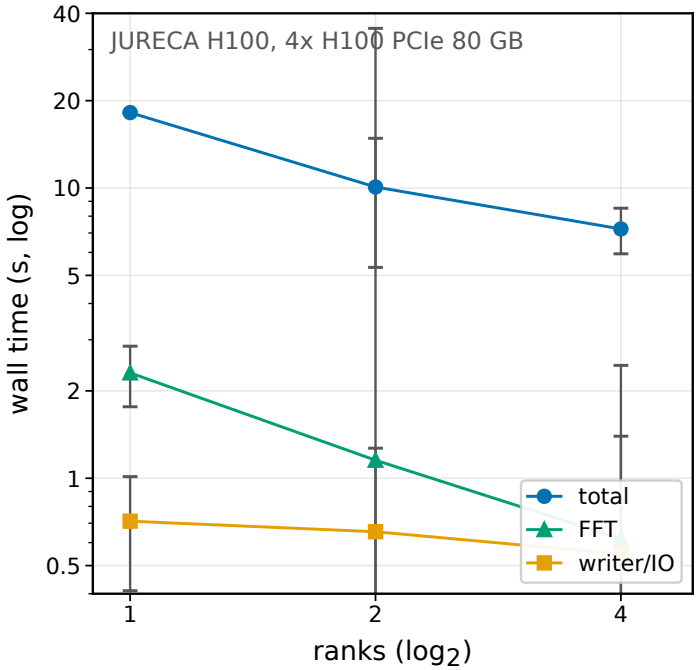


Figure 5: Distributed-mode strong scaling on JURECA H100 (4× H100 PCIe 80 GB) for 1–4 ranks, medians of three serialized runs: total, distributed FFT and distributed ADIOS2/BP5 writer at 256^3, $N = 2 \times 10^6$. The FFT reaches 3.73× and the total 2.52×; the writer scales sub-linearly, to 1.30× at four ranks.

many-rank bottleneck, not the field calculation.

The near shell still partitions tiles but all-reduces the full correction array; at 32^3 that communication and the tile imbalance dominate, so its rank scaling is not reported as a speedup. A local-slab correction reduction is required before the near shell is a many-rank feature.

Table 12: H100 measurement, medians of three serialized runs. Top: $N = 2 \times 10^6$, one bin, one rank, writer-inclusive stage times in ms, ADIOS2-BP5 backend (the V100 rows of Table 10 used HDF5; the 20-cell row's read-out is bimodal, see the text; "kernel" is the CLI kernel stage, "FFT" the CLI fft stage, distinct from the transpose).

box	deposit	kernel	FFT	read-out	writer	total
20 cells	36.5	1.3	0.51	23.4	5.47	67.50
100 cells	39.7	4.1	32.3	26.8	24.05	127.66

7.2 H100 measurement

A campaign on `jureca-h100` (four H100 PCIe 80 GB, sm_90, CUDA 13.4, serialized on one GPU, medians of three runs) reproduces the scaling picture with this hardware. The writer-inclusive rows below use the ADIOS2-BP5 backend (the V100 rows of Table 10 used HDF5): at $N = 2 \times 10^6$, a 20-cell box takes 67.50 ms end to end (deposit 36.5, kernel 1.3, FFT 0.51, read-out 23.4, writer 5.47 ms; writer share 8.1%); a 100-cell box takes 127.66 ms (deposit 39.7, FFT 32.3, read-out 26.8, writer 24.05 ms; writer share 18.8%, against 53.8 ms for the HDF5 writer). The deposit scales with N as in Table 9 (sub-linear below $\sim 10^6$, approaching linear at 10^6 and above); the 100-cell case is FFT- and writer-bound, and the ADIOS2-BP5 backend accounts for the whole -13.7% difference of the 100-cell total against the HDF5 value 147.9 ms (the compute stages agree). The 20-cell total is $+11.6\%$ above the HDF5 value 60.5 ms; its read-out is bimodal across the three repeats (18.2/23.4/70.8 ms, the slow one being the known OMP-over-unified-memory read-out mode), while the deposit stays flat at 36.2–36.6 ms and the two fast-mode repeats (62.0/67.5 ms) bracket the HDF5 value. For $N = 200$ particles the 16- and 64-bin multi-velocity runs take 92.2 and 100.6 ms; the B_z peak differs by 4.6%, within the shot noise of 200 particles.

Rank scaling at $N = 2 \times 10^6$ and the near-shell case are in Section 7.1: the BP5 writer scales sub-linearly to 1.30× at four ranks while the distributed FFT reaches 3.73×; the total (mesh stage plus writer) reaches 2.52×.

8 Discussion

The method computes the present-time field of a uniformly moving, TSC-shaped bunch as a free-space FFT convolution. It is a far-field method, and the measured error structure says exactly where that stands.

In the far field the relative error against the TSC-convolved continuum is below 3×10^{-4} on the beam axis from $d = 6$ for every γ; perpendicular to the beam the crossover is $d = 8$ for $\gamma = 1$, aspect 1, $d = 32$ for $\gamma = 1$, aspect 4, and $d = 48$ for $\gamma = 10$, aspect 1; the perpendicular value 24 in Table 4 for the same $\gamma = 1$, aspect 4 configuration is the crossover at the stricter 3×10^{-4} gate (perpendicular direction), a different threshold from the 10^{-3}-gate continuum crossover quoted here. The crossover distances are V100 values and were not re-measured in this campaign. The mesh against the point field carries the TSC smoothing in addition to the deposit error, and the two disagree by a factor of 58 in L2 at $\gamma = 10$, aspect 1, $d = 48$ (Table 2); a mesh-versus-point number measures the deposit shape, not the discretization. The cost is two regimes with a crossover set by the box, not the particle count: the deposit grows with N (sub-linear below $\sim 10^6$, approaching linear at the top of the sweep) and dominates at large N (33–38 ms on the 2×10^6 rows, one H100 of JURECA H100), while the FFT and read-out scale with M^3 and dominate at large boxes (51.9 to 85.7 ms from $M = 60$ to 120, and to the 100-cell row, which is bimodal with low/high modes 115.99/241.89 ms at $M = 220$). Repeated runs are bit-identical (md5 `d3641d64...`), and the Coulomb limit reproduces the TSC-convolved Coulomb field to 4.88×10^{-14}.

The near field is the limit of the mesh. Where the slice planes cut the bunch core, the mesh-only worst per-component plane error is 0.12–0.415 for $\gamma = 1$ (E_z) and 0.116–5.33 for $\gamma = 10$ (minimum E_y, x-y; maximum E_z, x-y) (Appendix C). The opt-in near shell assembles exactly to roundoff, and its tile subtraction is not the limiting term: at the measured $R = 3/4$ radii on the $h = 1$ slice grid the residual is the mesh contribution outside the shell, not the subtraction. The fast backend meets 10^{-3} stationary (3.71×10^{-4} interior) and moving axial (1.40×10^{-4} interior); the moving oblique interior is 1.94×10^{-3}, the measured limit of the fast backend at its default order (Table 6). The shell radius is the dominant cost lever: at 2×10^6 particles the support-only shell is 500.95 s end-to-end at 7.99×10^{-4} and the four-cell shell 1959.31 s at the same interior error, against a mesh-only wall of 0.500 s (Section 5.7).

The γ landscape states the mesh's sub-cell limit. At a fixed distance in cells the error grows steeply with γ: 0.26, 2.3, 5.3 for $\gamma = 10, 50, 100$ at $d = 8$, $h = 1$ (Table 5, Fig. 3). The controlling length is the beaming cone width d/γ; for $\gamma \geq 50$ the cone is sub-cell, and in the node-unit kernel the h-refinement exponent p decreases from 1.33 at $\gamma = 10$ to 1.09 at $\gamma = 100$ (three-point fit over $h = 4, 1, 1/4$ cells on the $d = 2$ band), so refinement continues to reduce the error across the range. Sub-cell resolution of the cone is the binding limit: at $d = 8$, $\gamma = 100$ the cone is 0.08 cells, and 10^{-2} requires h below that. At $\gamma = 50$ to 100 the near-charge error requires a shell of radius of order γ cells; that is the regime where the near shell is the working tool rather than the mesh.

The scope of the validated measurements is a single shared velocity, or a set of binned velocity classes with one convolution per class; the near shell needs no binning because each pair carries its own velocity. The measured binned behaviour agrees to 6.9×10^{-6} in the B_z peak and 1.1×10^{-3} in the small B_y component for the 5% energy spread. The 1 mrad case is within 1.4% of the arbiter in its original full-grid run and is not re-measured here (Appendix D). The computed quantity is the TSC-convolved uniformly-moving point-charge field; the convolution is free-space (no boundary condition other than zero fill beyond the far-tail switch); and there is no radiation term, because the motion is uniform. The checkpoint arrives bitwise under a stock PIConGPU, periodic and open-boundary, single- and multi-patch, so the field is usable as the start state of the simulation it is meant to initialize.

9 Conclusion

The method computes the present-time field of a uniformly moving, TSC-shaped bunch as a zero-padded free-space FFT convolution with staggered Yee read-out, driven by a single TOML configuration, and writes a uniform float32 openPMD checkpoint that restarts under a stock PIConGPU, periodic and open-boundary, single- and multi-patch. The implementation is validated on V100 `sm_70`, H100 `sm_90`, and gfx1100 HIP hardware.

The method is a far-field tool. In the far field the mesh is below 3×10^{-4} of the TSC-convolved continuum from $d = 8$ cells for $\gamma = 1$, aspect 1, and from $d = 32$ for $\gamma = 1$, aspect 4; at $\gamma = 10$, aspect 1, the measured values are 6.3×10^{-3} at $d = 32$ and 2.6×10^{-4} at $d = 48$ (Table 2, L2; the $d = 48$ entry is a V100 value not re-measured here). At 2×10^6 particles the mesh-only wall is 0.500 s (the mesh stage of the $R = 1.5$ cost row, compute 0.083 s) on one H100 of JURECA H100, deterministic (md5 `d3641d64...`), and the cost grows with the particle count for the deposit (sub-linear below $\sim 10^6$) and is cubic in the padded grid for the FFT. The opt-in near shell assembles exactly to roundoff and its fast backend meets 10^{-3} stationary and moving axial, with the moving oblique interior at 1.94×10^{-3}, the measured limit. The MPI compute is a component-parallel mesh and a tile-parallel near shell, bitwise identical at 1/2/4 ranks on JURECA H100 (four H100 PCIe), with the distributed FFT reaching a 3.73× speedup and the total (mesh stage plus writer) a 2.52× speedup at four ranks; the BP5 writer scales sub-linearly, to 1.30× from one to four ranks. The measured sub-cell limit at $\gamma \geq 50$ and the mild decrease of the fitted h-refinement order with γ are stated in Section 5.5; in that regime the near shell is the working tool.

Methods

9.1 Bunch and grid configurations

Table 13 lists every measured case with the bunch and the grid, and the section that uses it. The entries are the recorded configuration values of this work. The generation line for the mid-bunch bunches is gen erate_bunch.py--modelgamma--gammaG--seed42--grid646432--cell-size1e-71e-71e-7--centroid323216--si gma666--weightsuniform. The total charge is N times $-e$ for the electron bunches with one unit weight, where $e = 1.602176634 \times 10^{-19}$ C. The multi-velocity cases of Appendix D and the forward/fluctuation slice cases of Appendix C carry the macro weight 5×10^{-13} C per particle ($Q = -100$ pC for $N = 200$); all other cases are unit weight. The kinetic energy is $(\gamma - 1)\, m_e c^2$ with $m_e c^2 = 510.999$ keV. The mid-bunch bunches are mono-energetic: the read-back records $\gamma_{\min} = \gamma_{\max} = 1$ for the $\gamma = 1$ bunch and $\gamma_{\min} = \gamma_{\max} = 10$ for the $\gamma = 10$ bunch.

Table 13: Configurations of every measured case and the section that uses it. The two mid-slice entries give the bunch generation grid; the field grid for those cases ($32 \times 32 \times 16$ @ 2×10^{-7} m) is in the second block. Kinetic energy is $(\gamma - 1)m_e c^2$; the multi-velocity driver cell size is 10^{-6} m.

case	N	total charge (per-particle weight)	size (cells/axis)	spread	used in
mid slices, g10	1000	1.602×10^{-16} C	bunch generation grid $64{\times}64{\times}32$ @ 10^{-7} m ($6.4 \times 6.4 \times 3.2$), $\sigma{=}6$ cells/axis	none ($\gamma_{\min} = \gamma_{\max} =$ 10)	§C, §6
mid slices, g1	1000	1.602×10^{-16} C	bunch generation grid $64{\times}64{\times}32$ @ 10^{-7} m ($6.4 \times 6.4 \times 3.2$), $\sigma{=}6$ cells/axis	none ($\gamma_{\min} = \gamma_{\max} =$ 1)	§C, §6
large-domain, g10	100	1.602×10^{-17} C	extent $2.74 \times 3.78 \times 2.65$	...	§C (Fig. 16)
large-domain, g1	100	1.602×10^{-17} C	extent $2.74 \times 3.78 \times 2.65$	...	§C (Fig. 17)
cost, stationary	2×10^6	...	20-cell bunch	none	§5.7
cost, moving axial	2×10^6	...	20-cell bunch	none	§5.7
landscape, aspect 1	1	1.602×10^{-19} C	single particle, 96-cell box	none	§5.5
landscape, aspect 4	1	1.602×10^{-19} C	single particle at the origin	none	Table 2
multi-velocity, E 5%	200	1.0×10^{-10} C (-100 pC, weight $3.12 \times 10^6\, e$ each)	$\sigma = 6$ cells, centroid $(32, 32, 32)$	5% energy, $+y$ drift	§D
multi-velocity, A 5°	200	1.0×10^{-10} C (-100 pC, weight $3.12 \times 10^6\, e$ each)	$\sigma = 6$ cells, centroid $(32, 32, 32)$	5° angular, $+y$ drift	§D
1 mrad, binned	200	1.0×10^{-10} C (-100 pC, weight $3.12 \times 10^6\, e$ each)	$\sigma = 6$ cells, centroid $(32, 32, 32)$, 1.046 mrad rms x-div	1 mrad, $+y$ drift	§D, §C (Fig. 14)
divergence, 1 mrad	1000	1.602×10^{-16} C	$\sigma = 0.6$ each axis, centroid $(3.197, 3.208, 1.615)$	1 mrad paraxial, $+y$ (seed 43)	§C (Fig. 15; first 200 at unit weight, -32 aC)

case	γ, kinetic E ($m_e c^2 =$ 510.999 keV)	weight, cell	field grid (domain), host
mid slices, both	10 / 5.06 MeV; 1 / 0	1 each; 2×10^{-7} m	$32 \times 32 \times 16$ ($6.4 \times 6.4 \times 3.2$), rosi_v100
large-domain, both	10 / 5.06 MeV; 1 / 0	...; 2×10^{-7} m	$128 \times 128 \times 64$ ($25.6 \times 25.6 \times 12.8$), jureca-h100
cost, both	1 / 0; 10 / 5.06 MeV	...; 2.5×10^{-7} m	$66 \times 68 \times 66$, jureca-h100
landscape, aspect 1	10/50/100	1; 10^{-7} m	96^3 ($h{=}1$), 384^3 ($h{=}1/4$), jureca-h100 (ref: rosi_v100)
landscape, aspect 4	1/10/100	1; anisotropic, aspect 4	axis-aligned, $h{=}1$, rosi_v100
multi-velocity, both	10 / 5.06 MeV	$3.12 \times 10^6\, e$ (5×10^{-13} C) each; 10^{-7} m	64^3, JURECA H100
divergence	10 / 5.06 MeV	1 each; 2×10^{-7} m	$32 \times 32 \times 16$ ($6.4 \times 6.4 \times 3.2$), rosi_v100

9.2 Command and software versions

The binary is invoked as picBunch--inputbunch.h5--destination<checkpoint>--configparams.toml, with every unit and conversion parameter in the TOML file. The `[picongpu]` section carries `unit_length`, `unit_e`, `unit_b`, `dt`, `step`, `precision`, `species`; `[domain]` carries `cells`, `cell_size`, `origin`; `[placement]` carries `offset_unit` and the `left`/`front`/`top` fractions; `[accuracy]` carries `tolerance` and the near-shell settings (`near_shell none`/`exact`/`fast`/ `fast-gpu`, plus `near_shell_radius` and `near_shell_tile`); `patch_count` in `[placement]` splits the species into one particle patch per output rank for multi-rank restarts. A missing conversion key causes a clear error at start-up. The input bunch is centred by its charge-weighted centroid and placed at the requested offset from the domain origin; a placement leaving the bunch outside the domain is rejected at start-up.

The measurement hosts and toolchains are: `rosi_v100` (node `gv002`, V100 `sm_70`, CUDA 12.6, `PICBUNCH_ARCH=70`, Release); JURECA H100 (H100 `sm_90`, CUDA 13.4); the multi-velocity measurements of Appendix D, the hires and line series ran on JURECA H100 (H100 `sm_90`, CUDA 13.4); `hal` is used for the AMD Radeon RX 7900 XTX `gfx1100` ROCm 7.2.0 toolchain (rocFFT) in the distributed-FFT HIP run and the CTest portability check. The checkpoint writer follows openPMD standard version 1.1.0 [Huebl et al., 2018] and is built against the openPMD-api 0.17.1 software. The reference PIConGPU is PIConGPU 0.9.0-dev (deck commit `b87c053`), the release whose `referenceOnly` field-initializer mode is used in Section 6. The FFT library is cuFFT in the unnormalized form of the vendor documentation [NVIDIA Corporation, 2024].

Particle shape. The charge deposit uses the PIConGPU species shape selected with `-shape` (`cic`, `tsc` or `pqs`; default `tsc`). The support radii are 1, 1.5 and 2 PIC cells. The shape is not recorded in the openPMD particle data, so it must mirror the simulation's `shape.param`. The shape sets the source smoothing; the FFT kernel remains the sampled point kernel. The near shell supports TSC only and rejects CIC/PQS, while the mesh path supports all three.

Chunked particle input. `-max-particle-chunk-gb` (a float, for example 1.5) streams the selected species in chunks of at most that many gigabytes instead of loading it once. Each chunk is read on rank 0, broadcast, computed on the same fixed grid, and added to the output field. The mode requires explicit `-grid`, `-origin` and `-cell-size` and a single velocity class, and the global-domain writer is not applied in this mode. On a 200-electron test, 1-particle chunks reproduce the unchunked field to 4.8×10^{-16}.

9.3 Data and code availability

The source is available at `https://github.com/psychocoderHPC/picBunch`. The measurement data referenced by name in this paper are the files recorded in `docs/measurements/` of the repository: `matrix.csv` (the bench sweep behind Table 9), `w4b_rel_error.csv` (Fig. 2, Table 3), `coulomb.csv` (the Coulomb anchor), and the per-case `accuracy_results.json` and `n1b_padding_*.csv` files (Appendix B), with SHA-256 checksums in `MANIFEST.sha256`. The field-slice `.npz` arrays behind Figs. 13–15 and Figs. 16–17 are recorded in the measurement campaign (Appendix C); individual measurements name the artifact they used. The figures are rendered at their print width from the recorded data with the shared style module described in the repository documentation; each figure PDF in `figs/` is rendered from recorded inputs, and the figure audit records the minimum printed font size and colour-bar label of each.

A Derivations

A.1 The singular box

The in-support integral of Eq. (1) is split by writing the kernel as its divergence plus a smooth remainder. The divergence part is the eight-term asinh box

$$
\begin{aligned}
T(A,B,u) &= \int_0^A \int_0^B \frac{\mathrm{d}X\,\mathrm{d}Y}{\sqrt{u^2+X^2+Y^2}} \\
&= \begin{cases} A\,\mathrm{asinh}(B/A) + B\,\mathrm{asinh}(A/B), & |u| < 10^{-300}, \\ A\,\mathrm{asinh}\big(B/\sqrt{u^2+A^2}\big) + B\,\mathrm{asinh}\big(A/\sqrt{u^2+B^2}\big) & \\ \quad -u\,\arctan\big(AB/(u\sqrt{A^2+u^2+B^2})\big), & \text{otherwise.} \end{cases}
\end{aligned} \tag{13}
$$

with $A, B > 0$ (Eq. (11)). Each box face has 8 corners and 24 such calls per component, so the eight-term box is O(1) per pair. The form is even in u: u is the distance from the field point to the face plane, so a field point beyond a box face gives negative u, and the $|u| \to 0$ branch must fire on $|u|$, not u. The branch at $|u| < 10^{-300}$ is the $u \to 0$ limit, which is finite; it is not a regularization, it is the limit of the second branch as $|u| \to 0$.

The primitive is verified against a converged midpoint trapezoid reference: the double integral of the integrand over the unit box with $u = 1$ is -2.6348 for the converged midpoint trapezoid; the naive trapezoid gives -2.6333, 1.5×10^{-3} low, a trapezoid singularity artefact at the corner where the integrand is not smooth. The primitive is implemented in the fast evaluator header `include/picBunch/detail/FastNear.hpp` and the device reference `tools/closed-form/gpu_extract.cu`.

A.2 The smooth remainder

The remainder integrates $(S - W^*)K$ over the TSC knot panels, where S is the step function of the kernel's support and W^* is the TSC shape subtracted so that the integrand is smooth inside the support. Each axis is split at the TSC knots $(-1.5, -0.5, 0.5, 1.5)$ and at the anchor point s^*, the projection of the source onto the axis; on each panel a plain tensor Gauss–Legendre quadrature is applied. The default order is 5 per axis, 8000 nodes per in-support pair. For a moving bunch the box is γ-stretched for the axial direction: the beaming cone compresses the field into a pancake of width d/γ, and the stretched box matches that compression so the singular part stays O(1). The oblique direction is integrated directly in the Lorentz-contracted $d = MR$ frame, where the Liénard kernel is the isotropic Coulomb kernel: the support box maps to a parallelepiped and the TSC shape to a degree-6 piecewise polynomial along each ray, evaluated by a $16(\mu) \times 32(\phi) \times 4$(radial) rule whose angular part is polar-aligned with $\boldsymbol{\beta}$, with no remainder on that branch.

A.3 The h^3 Riemann factor

The sampled convolution evaluates the continuum integral of Eq. (1) by a Riemann sum whose nodes are spaced h cells apart, so each sample covers a physical volume h^3 (in cell units). The unnormalized FFT round trip scales by M^3, so the read-out applies the factor $h_x h_y h_z/(M_x M_y M_z)$ to convert the sum to the integral. This is the only place where the mesh spacing enters the amplitude; the scale-invariance of Section 3.6 follows from it, together with the homogeneity of degree -2 of the kernel and the cell-unit Yee offsets.

B Seven-velocity single-particle validation

Figures 6–12 show the mesh field for one electron on the final 128^3 grid, cell 5×10^{-8} m over the same $6.4\,\mu$m domain, against the independent order-12 continuum reference on the mid-z x-y plane. The seven velocities are $\gamma = 1.0$, 1.1, 1.5, 2, 10, 50 and 100. Each figure is one velocity case, with y horizontal and increasing left to right, and its rows are labelled with γ, β and the particle count $N = 1$. Table 15 lists the measured worst relative errors over the plane with the 5×5 nodes around the source excluded: 3.779×10^{-3}, 3.747×10^{-3}, 2.578×10^{-3}, 7.078×10^{-3}, 3.505×10^{-1}, 1.590 and 3.490 in velocity order. The corresponding full-plane values that include the source-coincident node are 7.010×10^{-2}, 6.745×10^{-2}, 7.726×10^{-2}, 9.569×10^{-2}, 0.6740, 7.095 and 4.263; that node is a singular point of the legacy reference (Section 5), so the full-plane column bounds the old reference, not the accuracy of the mesh. The error grows with γ because the field is compressed into a fraction of a cell: at $\gamma = 50$ and $\gamma = 100$ the Lorentz pancake is thinner than the mesh spacing, so those cases show the geometric sub-cell limit of the mesh, not a solver error. Raising the resolution from 64^3 to 128^3 at the same physical domain left the relative errors unchanged, so the growth with γ is the field compression rather than a discretization that finer cells remove. The few-times-10^{-3} floor at low γ is the reference/quadrature floor at the settings used, not the solver's limit. The peak squared fields scale as $1/\text{cell}^2$ compared with the 64^3 series (a factor 16 for the halved cell), a consistency check on the two campaigns.

Table 15: Seven-velocity single-particle validation on the final 128^3 grid, cell 5×10^{-8} m, 6.4 μm domain, against the order-12 continuum reference. Worst is the maximum relative error over the six components on the mid-z x-y plane with the 5×5 source neighbourhood excluded; the last column includes the source-coincident node. The committed `indfield.py` reference is not used at that node.

γ	β	worst, 5×5 excluded	full plane
1.0	0	3.779×10^{-3}	7.010×10^{-2}
1.1	0.4166	3.747×10^{-3}	6.745×10^{-2}
1.5	0.7454	2.578×10^{-3}	7.726×10^{-2}
2.0	0.8660	7.078×10^{-3}	9.569×10^{-2}
10	0.9950	3.505×10^{-1}	0.6740
50	0.99980	1.590	7.095
100	0.99995	3.490	4.263

Padding and R2C/C2R details

With pad = 1 the cyclic convolution wraps and a 9.275×10^{-3} periodic-image artefact appears for a single stationary particle in a 20^3 box. For a single stationary particle the pad = 2 result agrees with the direct non-cyclic sum to 3.811×10^{-16} and the pad = 3 result to 1.905×10^{-16}; at pad ≥ 2 the error is at machine precision ($\leq 10^{-15}$). The artefact grows for beamed sources: at pad = 1 the 100-source $\gamma = 10$ cloud shows a 7.5×10^{-1} worst-component error.

For the R2C/C2R half-spectra the DC bin and the seven self-conjugate even Nyquist corners must be exactly real; the intermediate Nyquist planes are Hermitian-paired, not real. The R2C/C2R path matches the C2C path to 1.665×10^{-16} over all six components.

The γ-landscape runs of Section 5.5 use padding 1; the $d = 48$ rows carry a small periodic-wraparound floor of order 10^{-4} to 10^{-3}, so the $d = 48$ values are not padding-clean.

C Mid-bunch slices and large-domain slices

C.1 Mid-bunch slices: field and error side by side

The distance sweeps of Section 5.4 sample the far field. To measure the mesh in the bunch core we slice through it and compare, pixel by pixel, with an exact TSC-convolved reference: 64×64 planes, cell 10^{-7} m, $N = 200$, $Q = -100$ pC, $\gamma = 10$ with momentum along $+y$. Two cases share the positions and differ only in a 1.046 mrad rms transverse divergence in x (Figs. 13 and 14). A separate unit-weight bunch (the first 200 of the 1000-particle seed-43 case of Table 13, $Q = -32$ aC) propagates along $+y$ and is different from the seed-42 mono-energetic bunch that moves along $+x$ in the $\gamma = 10$ case of Table 16. It is binned on the $32 \times 32 \times 16$ grid of cell 2×10^{-7} m and sliced at its mid-plane (Fig. 15). In every panel of these $+y$ figures, y is horizontal and increases left to right, so the bunch moves rightwards and y is the beam axis; the vertical axis is x in the x-y cut and z in the z-y cut. The top row of each plate is a charge-density overview of the bunch, not a field map; its three panels are the coordinate planes and its transverse panels carry the "beam axis (out of plane)" note (see the caption of Fig. 13). Each figure has a field row with four columns (E_x, E_y, E_z, E^2); the forward and 1 mrad figures carry an error row underneath, while the reference-only fluctuation figure shows the electric and magnetic field rows alone. The three components share one symmetric-logarithmic colour scale and bar, and each component panel states its peak as a percentage of the shared maximum. E^2 (or B^2) has its own logarithmic scale from 10^{-6} of the panel peak to the peak, with its own unit label. The error row is logarithmic, floored at 10^{-6} and marked at the 10^{-3} gate. Pixels are masked in light grey where the local field magnitude is below 10^{-3} of the largest component peak on the plane. The forward case compares the mesh plus fast shell against an order-4 adaptive Gauss–Legendre driver. The Lorentz-contracted in-support quadrature of Section 5.6 is the present reference; it differs from that driver at the few-percent level at the closest node, so the in-support error pixels in these panels bound the in-support error from above. The dark outline marks the TSC support. The quantitative in-support statement is the fast-versus-arbiter 1.6% at the worst node (Section 5.6). The fluctuation panels show the complete reference field; the binned momentum-spread case is carried through the multi-velocity path (Appendix D).

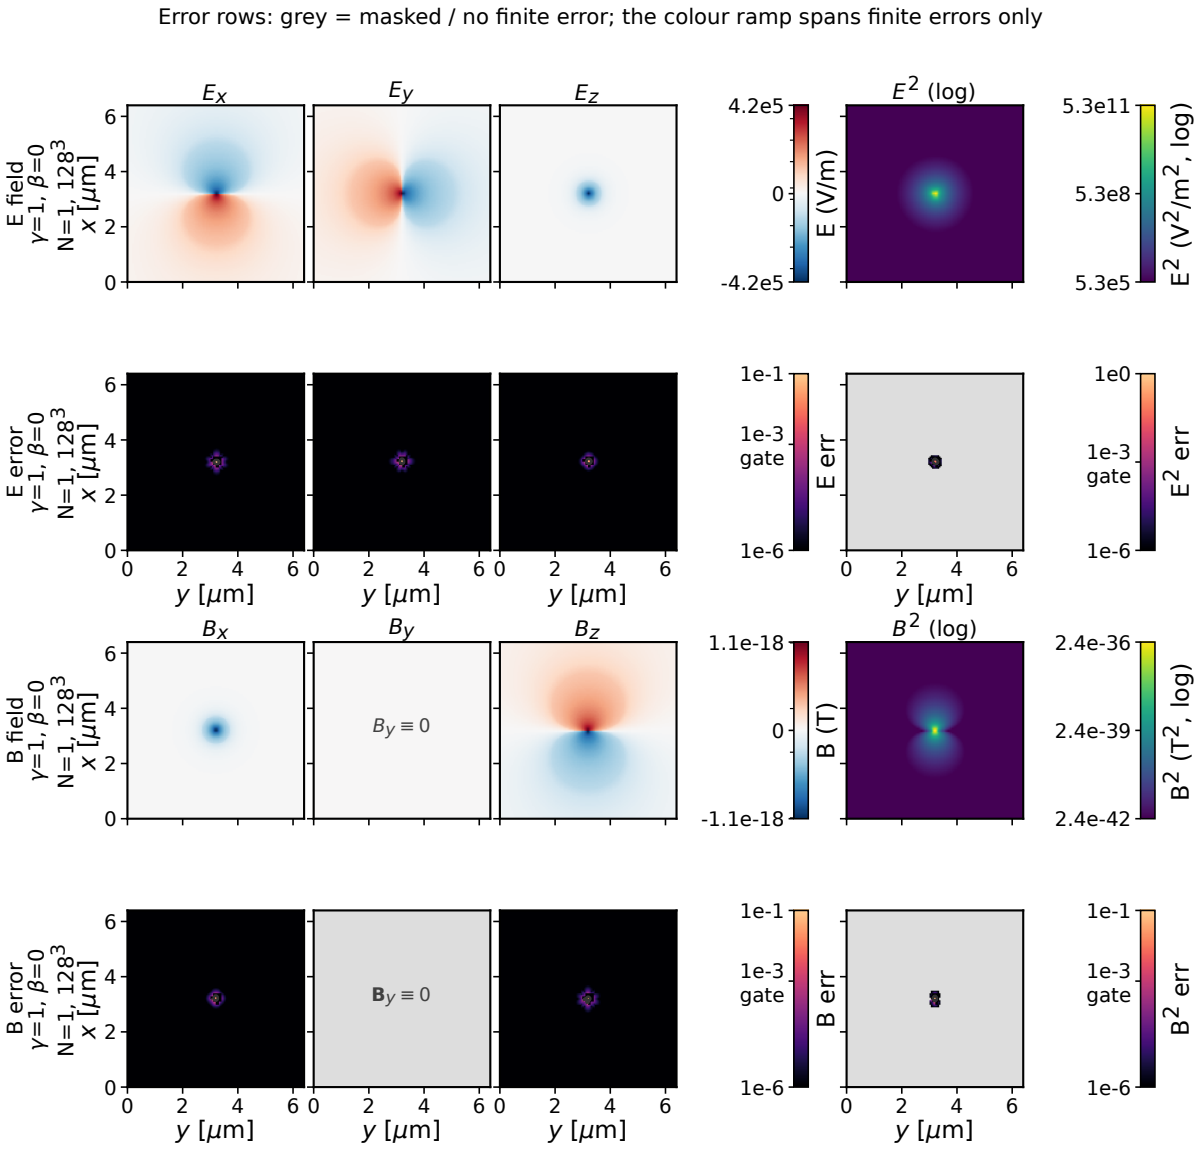


Figure 6: Single-electron field and error on the x-y plane at mid-z for $\beta = 0$ ($\gamma = 1.0$) on the 128^3 grid, cell 5×10^{-8} m. The $+y$ direction is horizontal and increases to the right, x is vertical; the bunch is stationary ($\beta = 0$). The E block and the B block each put E_x, E_y, E_z (or B_x, B_y, B_z) on one shared colour scale with its own bar and E^2 (or B^2) on a separate colour scale with its own bar, so the squared-magnitude scale never dominates the component scale. Error rows are logarithmic, carry the 10^{-3} contour. Grey pixels are masked: the reference field there is below 10^{-3} of the largest component peak on the plane, so the relative error is undefined. At $\gamma = 1.0$ the B error row is omitted entirely, since all three reference magnetic components (and B^2) are identically zero at $\beta = 0$ and the relative magnetic error is therefore undefined. The B field panels are blank for the same reason; the worst component error is 3.779×10^{-3} with the 5×5 source neighbourhood excluded (7.010×10^{-2} including the source node).

Shell radius unit. The near-shell envelope measured in the crossover work, "3 to 4 PIC cells," assumed a mesh spacing $h = 1/4$ of the PIC cell, so $R = 3$ and $R = 4$ are 12 and 16 mesh nodes. On this grid the node spacing is $h = 1$ PIC cell (2×10^{-7} m), so the same $R = 3/4$ is $3/4$ mesh nodes, four times smaller in physical units than the envelope configuration. At $R = 3$ on an $h = 1$ grid the shell covers 3 mesh nodes; the TSC support extends ± 1.5 cells from each source, so the shell covers a fraction of the full support. The measured plane errors at this radius are reported below.

The error is largest where these planes cut the bunch core: the spectral far-field convolution is not the dominant term there, and the maps show the measured error concentrated at the beamed core and falling away toward the plane edges. The far field is measured separately (Table 3, Fig. 2); the near-shell contribution is measured at the same geometry below.

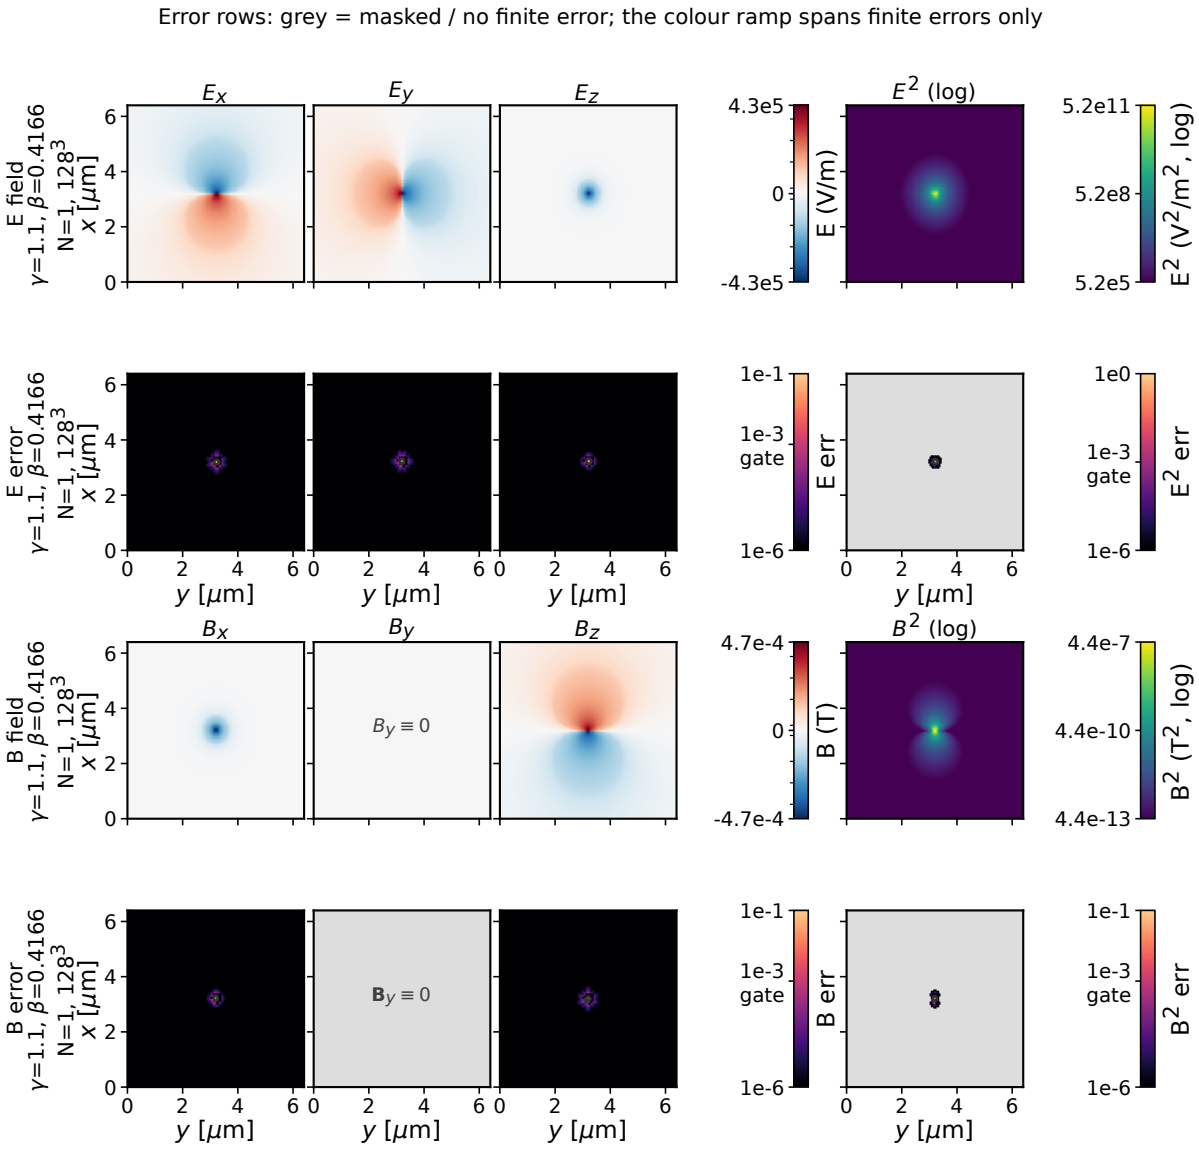


Figure 7: Single-electron field and error on the x-y plane at mid-z for $\beta = 0.4166$ ($\gamma = 1.1$) on the 128^3 grid, cell 5×10^{-8} m, with the same field, colour-bar and error conventions as Fig. 6 (grey pixels are masked). The B_y error panel is fully masked in every figure of this series, since $B_y \equiv 0$ for $+y$ motion and the relative error of a zero reference is undefined. The worst component error is 3.747×10^{-3} (6.745×10^{-2} including the source node).

C.2 Near-shell accuracy and cost on the slice geometry

On the same $N = 1000$, $32 \times 32 \times 16$, cell 2×10^{-7} m geometry (the seed-42 mono-energetic bunch of Table 16, beam axis $+x$; not the seed-43 $+y$ bunch of Fig. 15), the exact near shell was run at the radius the envelope case assumed, $R = 3$ for $\gamma = 1$ and $R = 4$ for $\gamma = 10$, in mesh nodes of this $h = 1$ grid; the mesh-only table and the near-shell numbers in this subsection are the v0.9.2 library runs, in contrast to the v0.4.0 standalone driver of the crossover work. The mesh-only plane errors (Table 16) are 0.12–0.415 for $\gamma = 1$ (E_z in both planes) and 0.116–5.33 for $\gamma = 10$ (minimum E_y, x-y; maximum E_z, x-y; beam axis $+x$, the B_x component identically zero): these planes cut the sources, the region where the spectral convolution is not the dominant term. With the near shell the worst per-component plane error is 1.7×10^{-3}–9.0×10^{-3} for $\gamma = 1$ (z-y plane E_x) and 6.7×10^{-2}–1.8×10^{-1} for $\gamma = 10$ (z-y plane E_x): at $R = 3/4$ on an $h = 1$ grid the shell radius covers $3/4$ mesh nodes, which is a fraction of the TSC support extent, and the measured residual is the mesh contribution outside the shell. The v0.4.0-driver high ends, 0.1067 for $\gamma = 1$ and 0.5704 for $\gamma = 10$, both sit on the same z-y E_z plane's last domain row ($z = 15$), where the v0.9.2 runs give 2.5×10^{-5} and 0.0954; the driver edge values are artifacts, not a

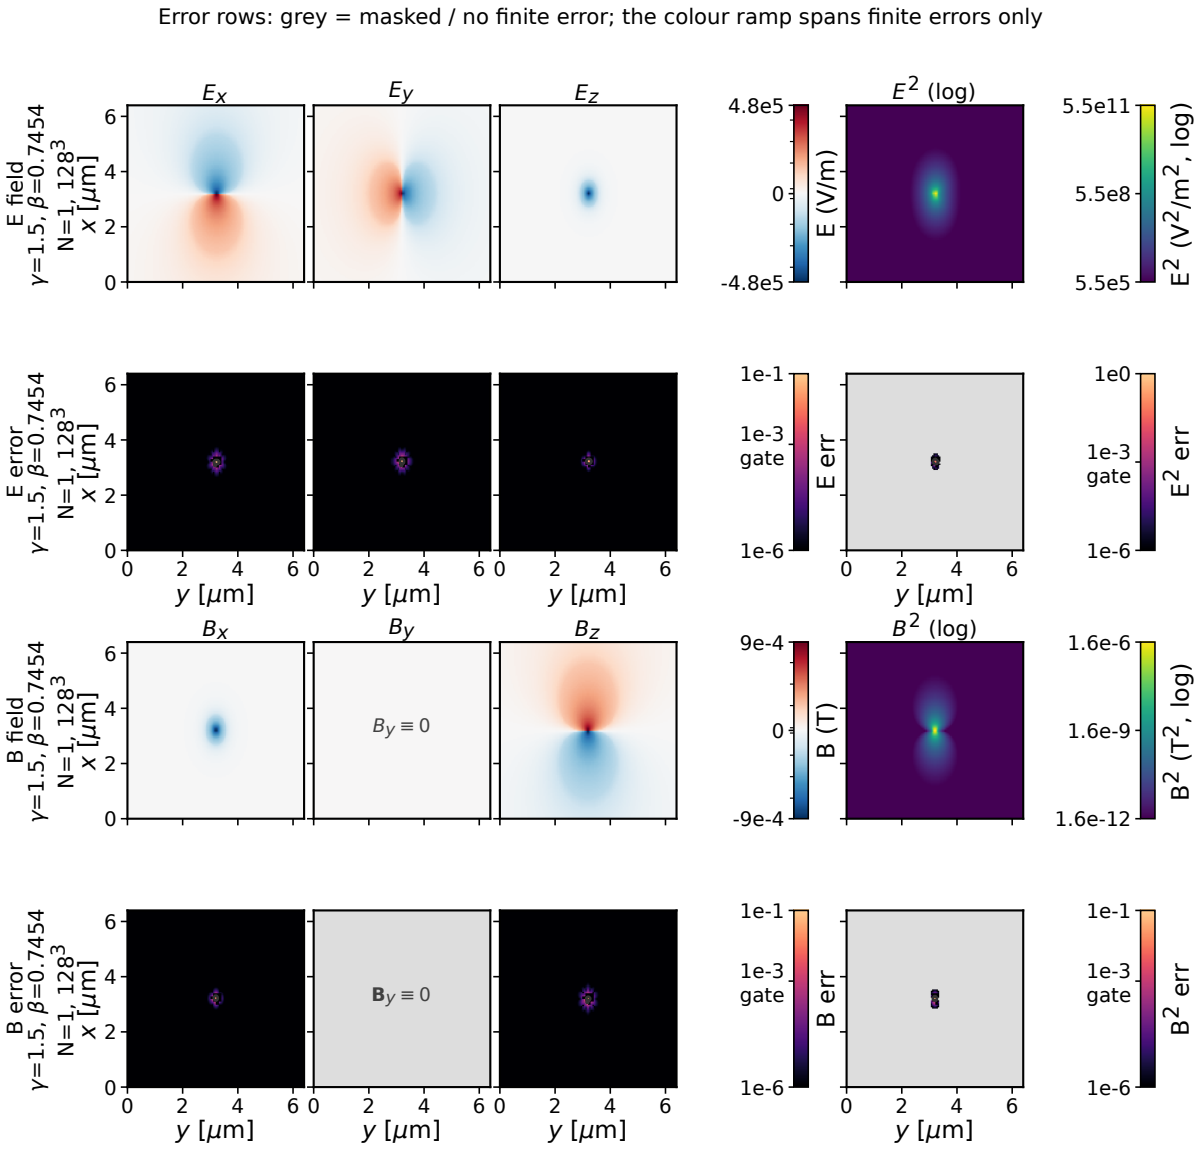


Figure 8: Single-electron field and error on the x-y plane at mid-z for $\beta = 0.7454$ ($\gamma = 1.5$) on the 128^3 grid, cell 5×10^{-8} m, with the same field, colour-bar and error conventions as Fig. 6 (grey pixels are masked). The worst component error is 2.578×10^{-3} (7.726×10^{-2} including the source node).

missing plane or radius. The discrete Gauss-law indicator, the Yee backward-difference div E residual normalised to the peak deposited ρ/ε_0, moves from 0.3534 to 0.0895 for $\gamma = 1$ and from 5.026 to 0.399 for $\gamma = 10$, against reference values 0.0904 and 0.0814: $\gamma = 1$ reaches the reference value, $\gamma = 10$ drops 12.6× and stays 4.9× above it.

The assembly itself is not the limiting term: the tests of Section 3.7 show the tiled subtraction exact to roundoff, so the residuals above are the mesh outside the shell. Runtime at $N = 1000$ is 4.3 min for $\gamma = 1$ and 32.7 min for $\gamma = 10$ (one 16-thread socket; the exact shell is the whole wall time); the exact evaluator costs about 8× more per pair at $\gamma = 10$ than at $\gamma = 1$.

C.3 Large-domain slices

Figure 16 shows the $\gamma = 10$ electric and magnetic field and error on a $128 \times 128 \times 64$ domain, cell 2×10^{-7} m, $N = 100$, against the independent continuum at tolerance 10^{-3}. The electric block comes first and the magnetic block second in the same figure, as in the per-γ presentation. Fig. 17 is the stationary-bunch case, with the electric block and the identically zero magnetic field row. Figs. 16–17 use the axial-x geometry: the bunch moves along $+x$, so B_x, the field component along the propagation axis, is identically zero. The x-y cut is longitudinal, with the beam axis x horizontal. The z-y cut is a

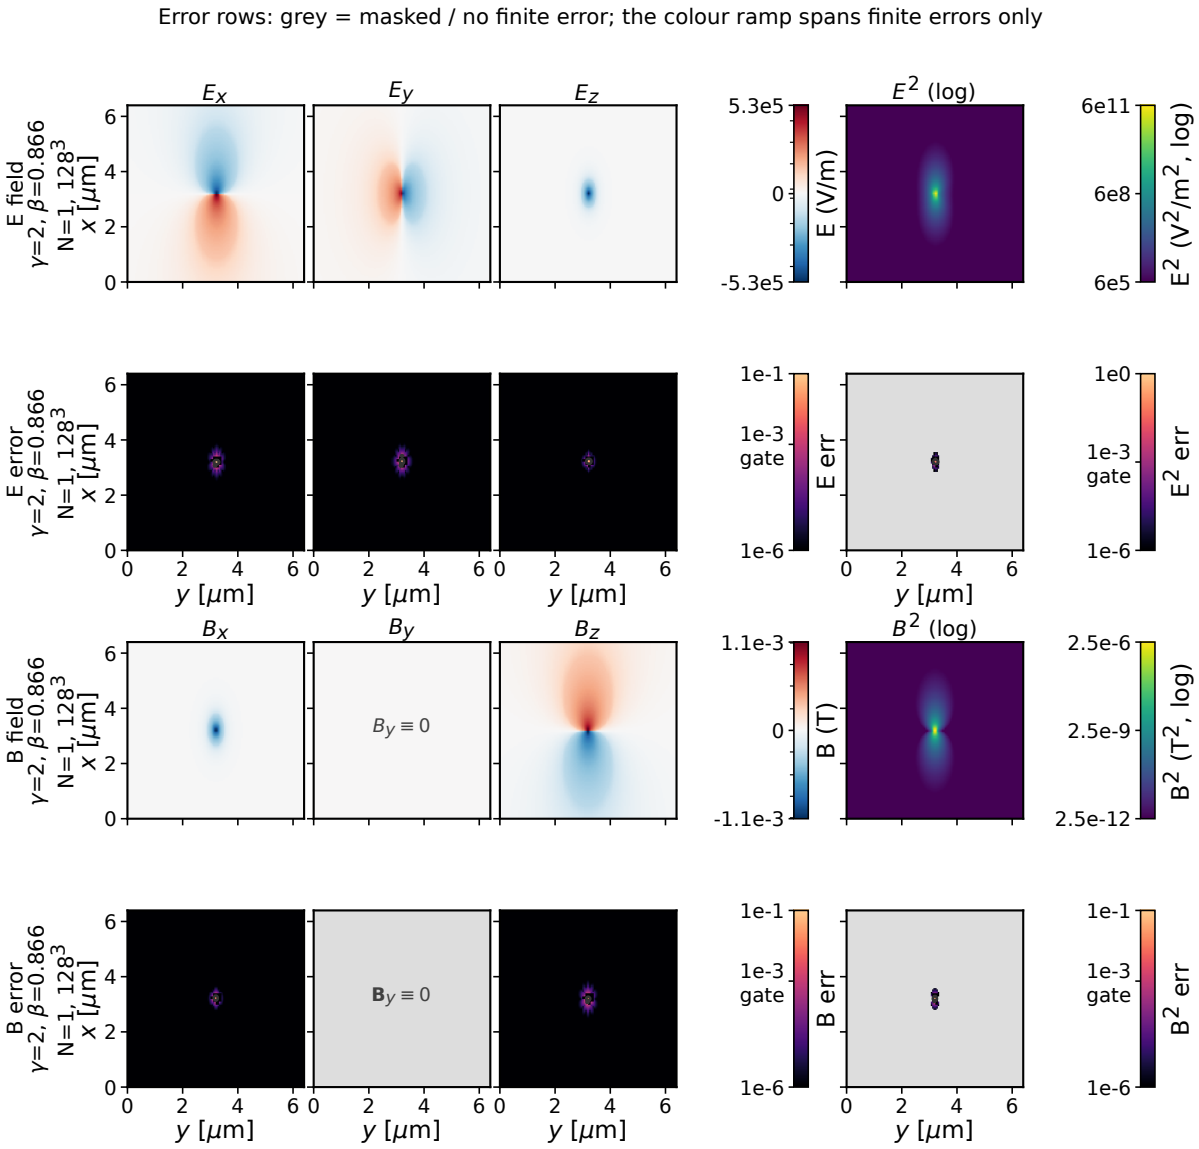


Figure 9: Single-electron field and error on the x-y plane at mid-z for $\beta = 0.8660$ ($\gamma = 2.0$) on the 128^3 grid, cell 5×10^{-8} m, with the same field, colour-bar and error conventions as Fig. 6 (grey pixels are masked). The worst component error is 7.078×10^{-3} (9.569×10^{-2} including the source node).

transverse view of the bunch: the beam axis x is perpendicular to the page, so the horizontal coordinate there is the transverse y, not the propagation direction, and z is vertical. This keeps the same horizontal y and vertical z as the z-y cuts of the $+y$ figures, where the out-of-plane transverse axis is x. The colour scale is set by the bunch, so the far decay is visible inside the frame; the core error is the method error of the mesh where the planes cut the bunch, far above the 10^{-3} gate.

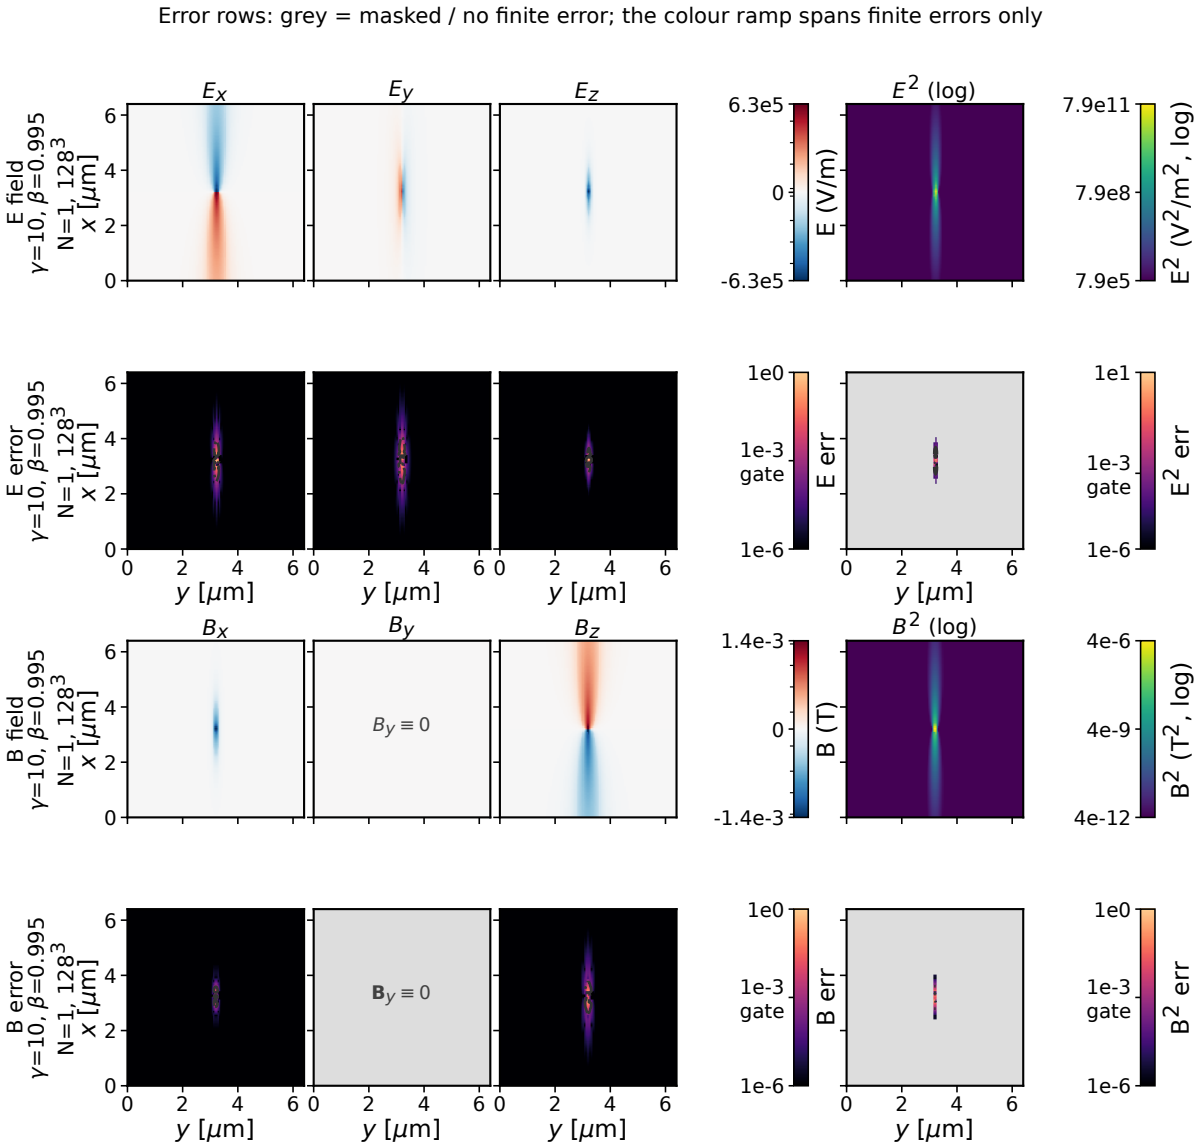


Figure 10: Single-electron field and error on the x-y plane at mid-z for $\beta = 0.9950$ ($\gamma = 10$) on the 128^3 grid, cell 5×10^{-8} m, with the same field, colour-bar and error conventions as Fig. 6 (grey pixels are masked). The worst component error is 3.505×10^{-1} (0.6740 including the source node).

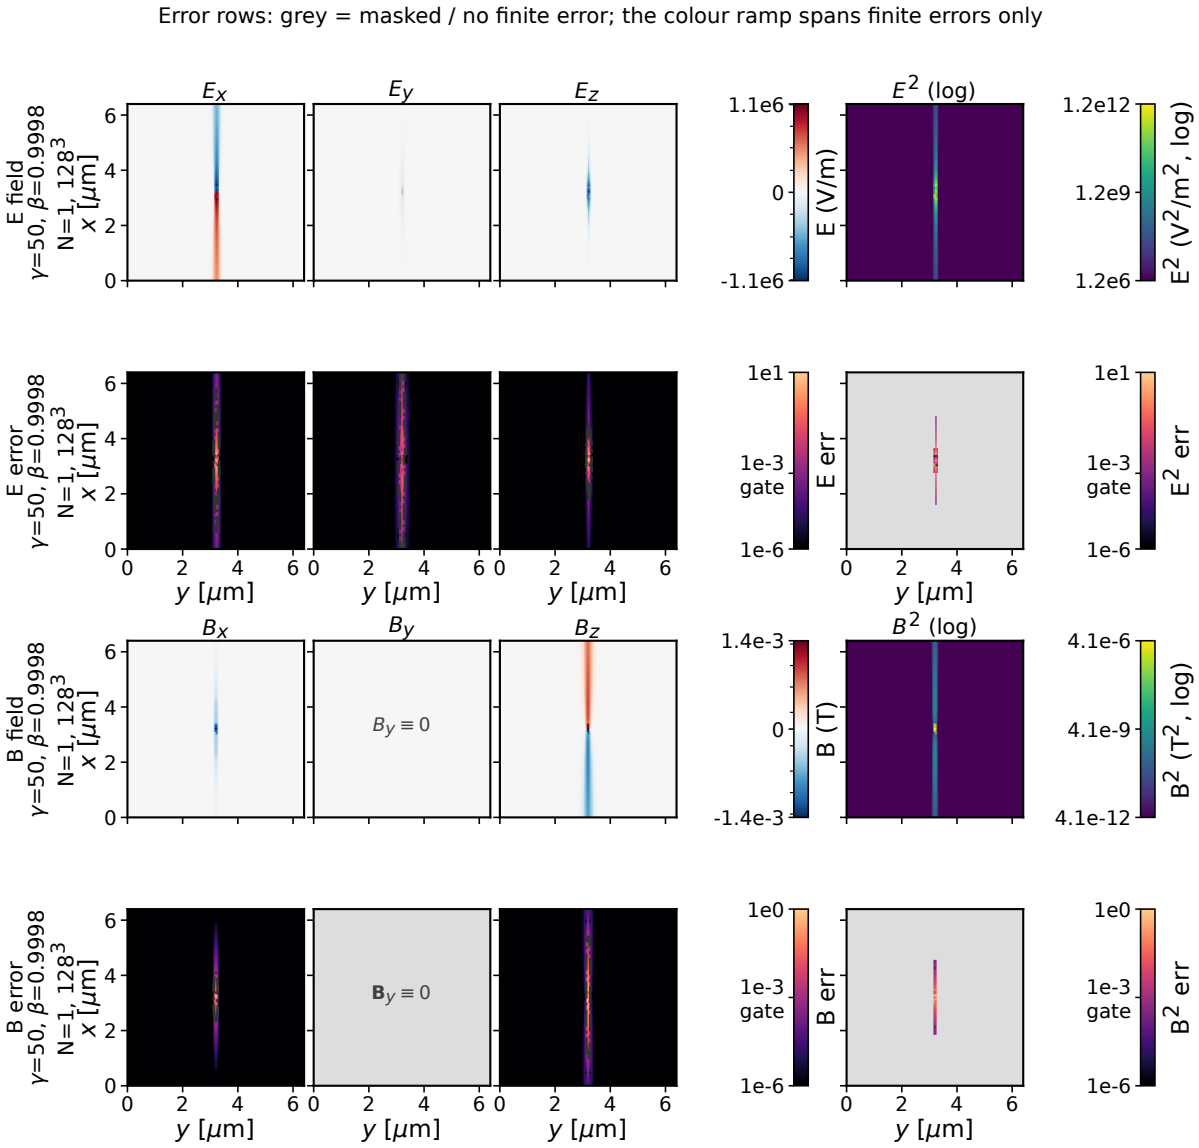


Figure 11: Single-electron field and error on the x-y plane at mid-z for $\beta = 0.99980$ ($\gamma = 50$) on the 128^3 grid, cell 5×10^{-8} m, with the same field, colour-bar and error conventions as Fig. 6 (grey pixels are masked). The worst component error is 1.590 (7.095 including the source node).

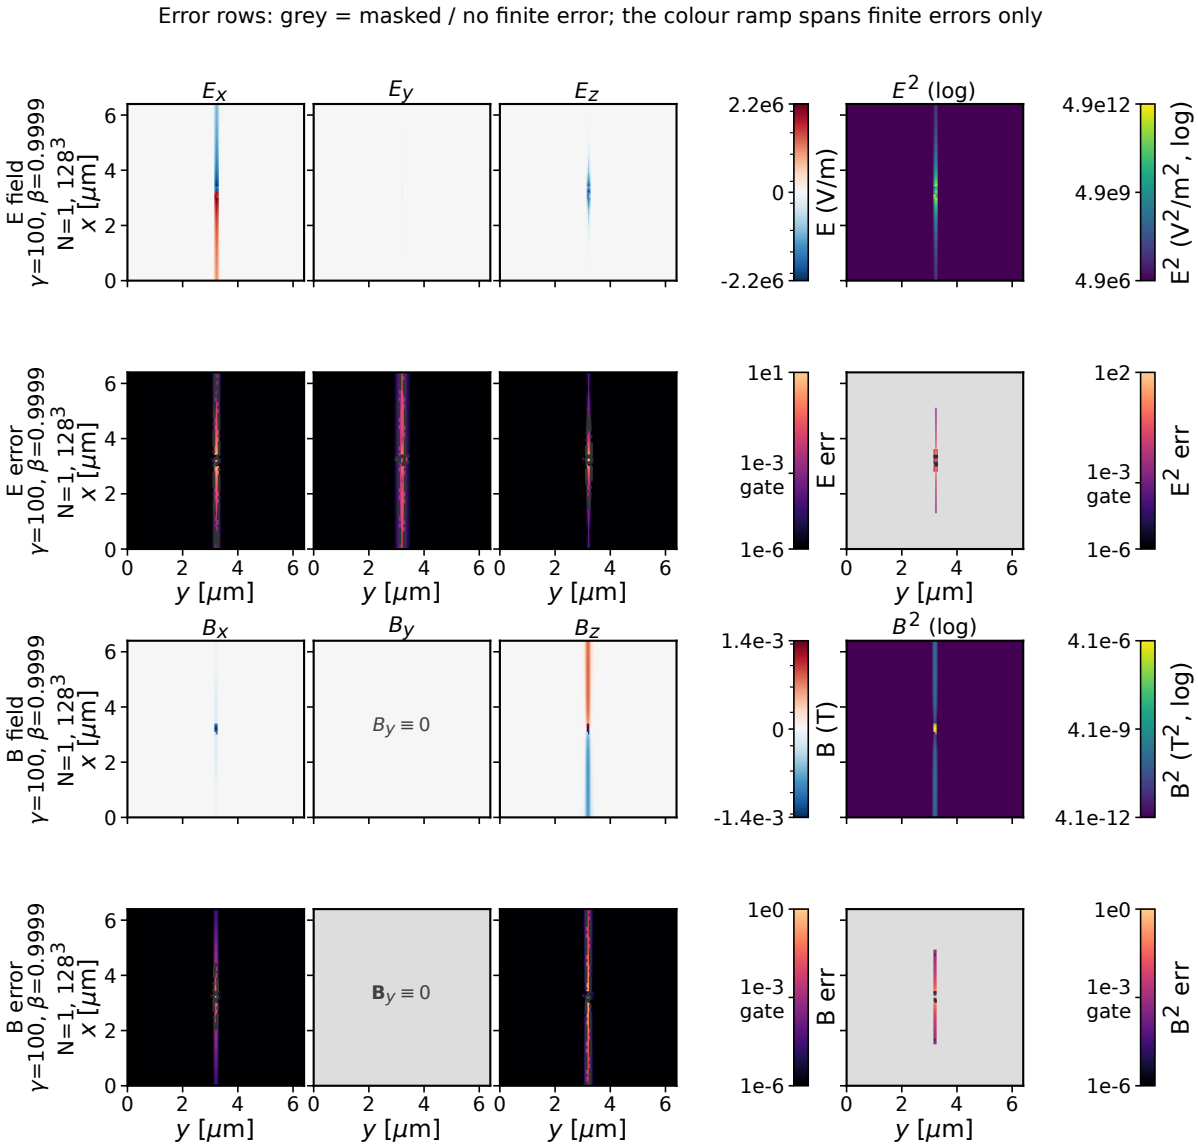


Figure 12: Single-electron field and error on the x-y plane at mid-z for $\beta = 0.99995$ ($\gamma = 100$) on the 128^3 grid, cell 5×10^{-8} m, with the same field, colour-bar and error conventions as Fig. 6 (grey pixels are masked). The worst component error is 3.490 (4.263 including the source node).

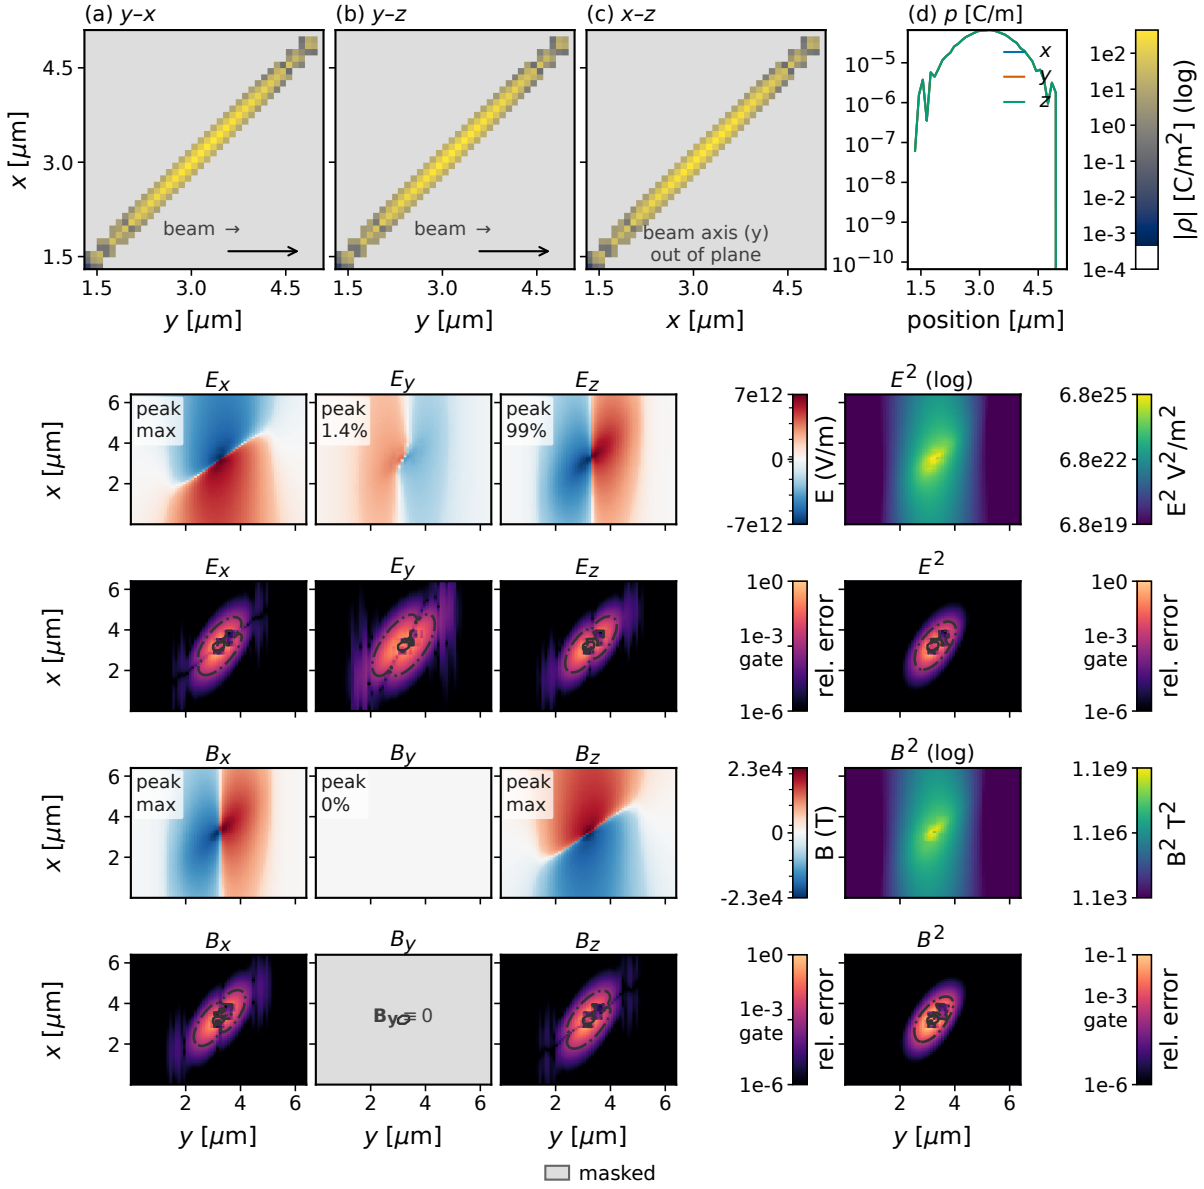


Figure 13: Forward-case validation ($N = 200$, $Q = -100\,\mathrm{pC}$, $\gamma = 10$ along $+y$, 64×64 planes, cell 10^{-7} m; x-y plane at $z = 3.2\,\mu\mathrm{m}$, with y horizontal, the beam axis; the z-y plane at $x = 3.2\,\mu\mathrm{m}$ is continued in Fig. 13. The top row is a compact charge-density overview of the same bunch: $|\rho|$ on a logarithmic scale floored at 10^{-6} of the peak, in the three coordinate planes on one shared colour bar, together with a marginal bunch profile. The `beam→` marker gives the propagation direction, and panels transverse to it are marked "beam axis (out of plane)". Below the density row, each image carries the electric block first and the magnetic block second. The electric block shows four columns E_x, E_y, E_z, E^2: the three components share one symmetric-logarithmic scale and bar (the text gives the E^2 scale and unit). Each component panel is labelled with its peak as a percentage of the largest component peak on that plane (the shared colour scale): the panel that sets the maximum reads `max`, peaks below one percent use one significant figure (e.g. `.04%`), and only an exactly zero peak reads 0%; a small percentage means the component is faint on that scale. The magnetic block shows the same layout for B_x, B_y, B_z, B^2. B_y is identically zero: its field panel is blank and its error panel is fully masked, since the relative error of an identically zero component is undefined. Under each field row lies a logarithmic error row, floored at 10^{-6} and marked at the 10^{-3} gate, with light-grey pixels where the local field magnitude is below 10^{-3} of the largest component peak on the plane; the dark outline marks the TSC support. The in-support fast-versus-arbiter error against the order-12 independent continuum reference is 1.6% at the worst node. The magnetic field is about 3.3×10^{-9} of the E-dominated local field at $\gamma = 10$ (Section 5).

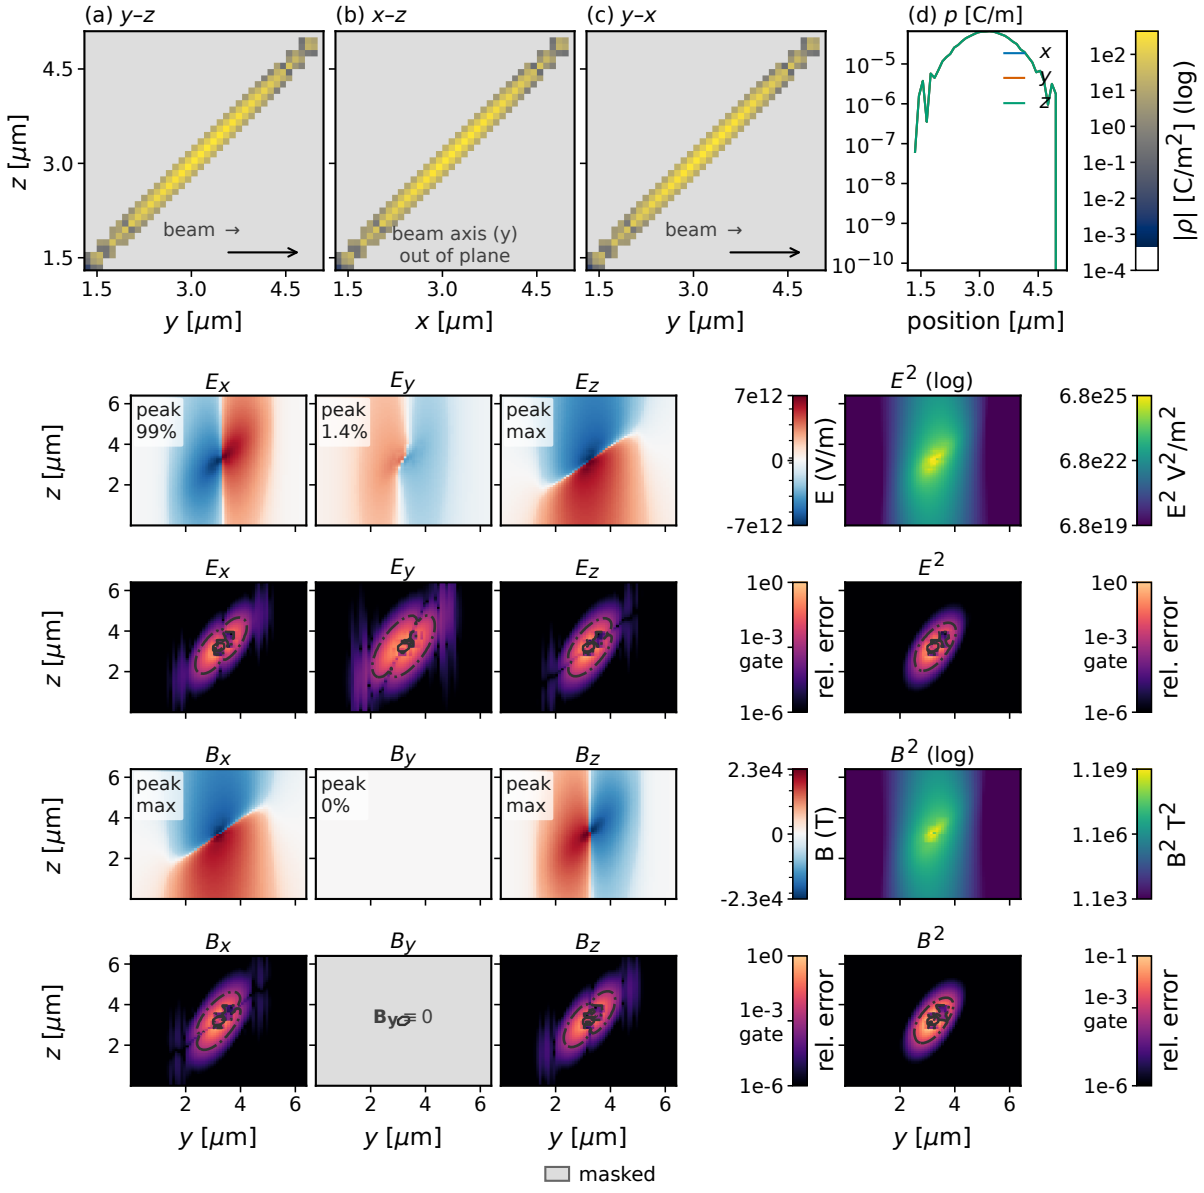


Figure 13: Forward-case validation (continued): z-y plane at $x = 3.2\,\mu\text{m}$, with the same electric/magnetic blocks, colour scales, masks and error conventions as the x-y plane in Fig. 13.

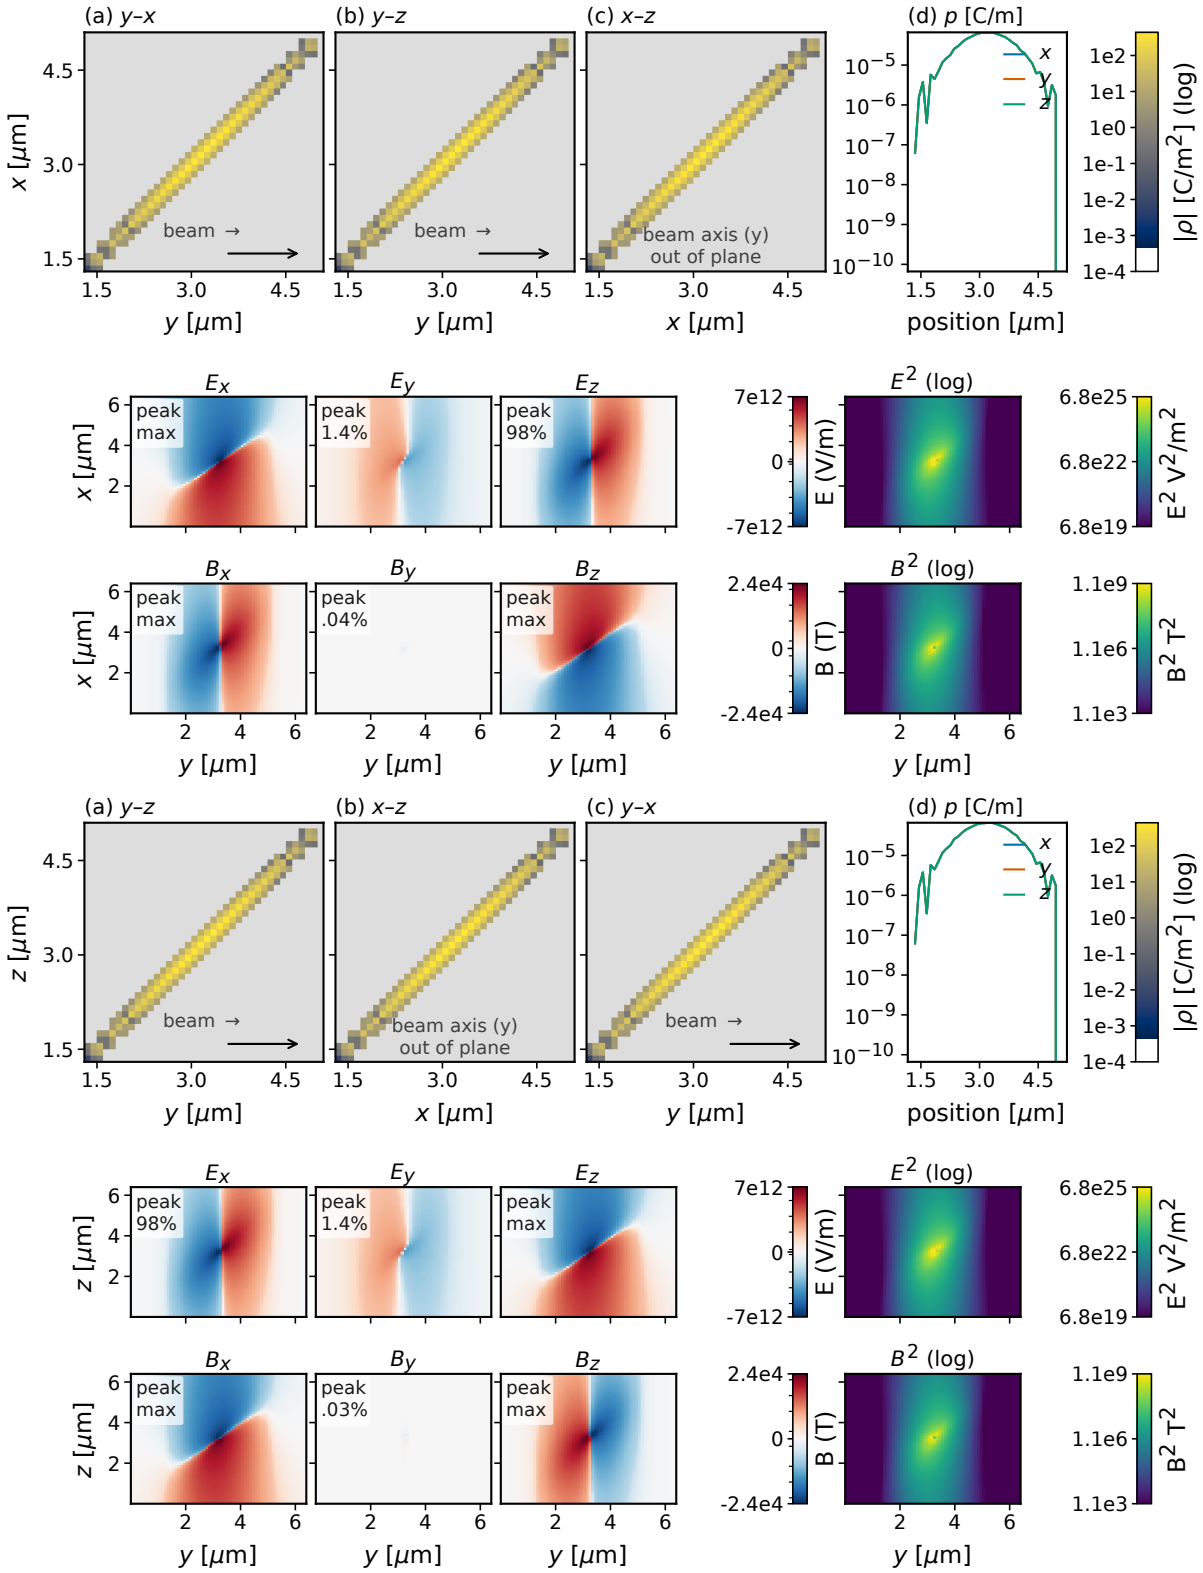


Figure 14: Fluctuation-case reference with the same geometry and bunch as Fig. 13 but a 1.046 mrad rms x-divergence. The top row repeats the charge-density overview of Fig. 13. The upper and lower images are the x-y and z-y planes, with y horizontal and the beam axis; each carries the electric field row first and the magnetic field row second, with the scales and masks of Fig. 13. The panels show the order-12 independent continuum reference; the x-divergence lifts B_y off zero, although it remains about 3×10^{-4} of the dominant component. The corresponding binned picBunch output is reported in Appendix D.

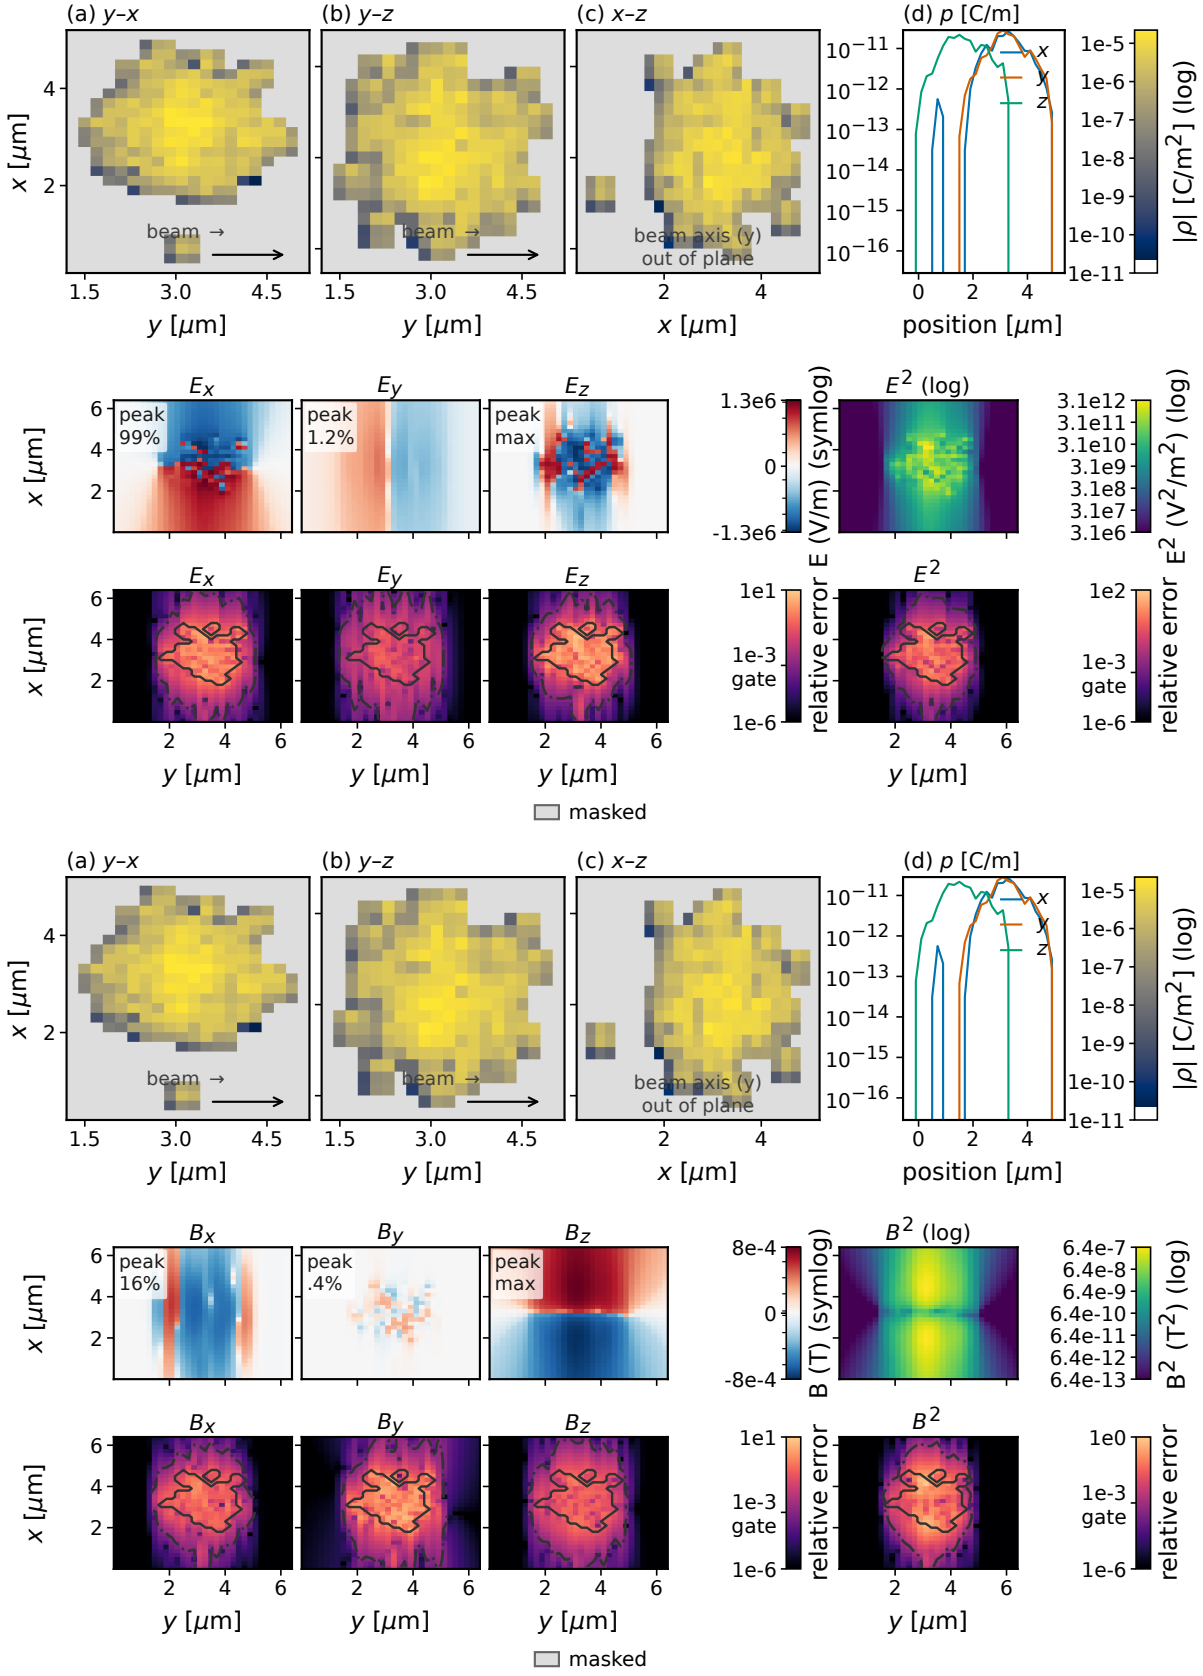


Figure 15: Binned 1 mrad case (seed-43 unit-weight bunch, $N = 200$, $Q = -32\,\mathrm{aC}$, $\gamma = 10$, 32×32 x-y plane at $z = 1.6\,\mu\mathrm{m}$, cell $2 \times 10^{-7}\,\mathrm{m}$), with the picBunch velocity-binned field against the order-12 independent continuum reference. The upper image is electric and the lower magnetic.

Table 16: Mid-bunch slice statistics, mesh-only. Worst relative error per component and fraction of unmasked plane pixels above 10^{-3} (unmasked means the local field magnitude is at least 10^{-3} of the largest component peak on the plane), x-y plane at $z = 1.60\,\mu\text{m}$ and z-y plane at $x = 3.20\,\mu\text{m}$, $N = 1000$, $32 \times 32 \times 16$, cell 2×10^{-7} m, mesh against the independent exact TSC driver (FP64 quadrature, order 4, max level 5, tolerance 10^{-3}). The $\gamma = 10$ case is the seed-42 mono-energetic bunch, which moves along $+x$ ($B_x \equiv 0$); x is horizontal in the x-y cut and is the beam axis there, while y is horizontal in the z-y cut (beam out of plane) and z vertical. Node spacing $h = 1$ PIC cell; no near shell. Measured on `rosi_v100` (V100, v0.9.2; the near-shell numbers in the text below this table are the JURECA H100 v0.9.2 runs, so this mesh-only table and those near-shell values are on different hosts). The $\gamma = 1$ reference is stationary (no beam axis) and uses the archived tolerance 10^{-4} (driver floor 9.075×10^{-4}, 14 not converged).

plane	component	worst rel. err	frac $> 10^{-3}$
$\gamma = 10$ *(beam axis $+x$; $B_x \equiv 0$)*			
x-y ($z = 1.60\,\mu\text{m}$)	E_x	1.81	54.1%
	E_y	1.16×10^{-1}	49.9%
	E_z	5.33	78.3%
	B_y	6.49×10^{-1}	72.7%
	B_z	3.77×10^{-1}	58.7%
z-y ($x = 3.20\,\mu\text{m}$)	E_x	5.06	90.4%
	E_y	1.17×10^{-1}	61.1%
	E_z	1.92	82.4%
	B_y	4.34×10^{-1}	80.9%
	B_z	7.16×10^{-1}	84.3%
$\gamma = 1$ *(stationary; all $B \equiv 0$)*			
x-y ($z = 1.60\,\mu\text{m}$)	E_x	1.20×10^{-1}	22.8%
	E_y	1.34×10^{-1}	23.3%
	E_z	4.15×10^{-1}	74.7%
z-y ($x = 3.20\,\mu\text{m}$)	E_x	4.15×10^{-1}	49.6%
	E_y	1.37×10^{-1}	61.7%
	E_z	1.23×10^{-1}	53.9%

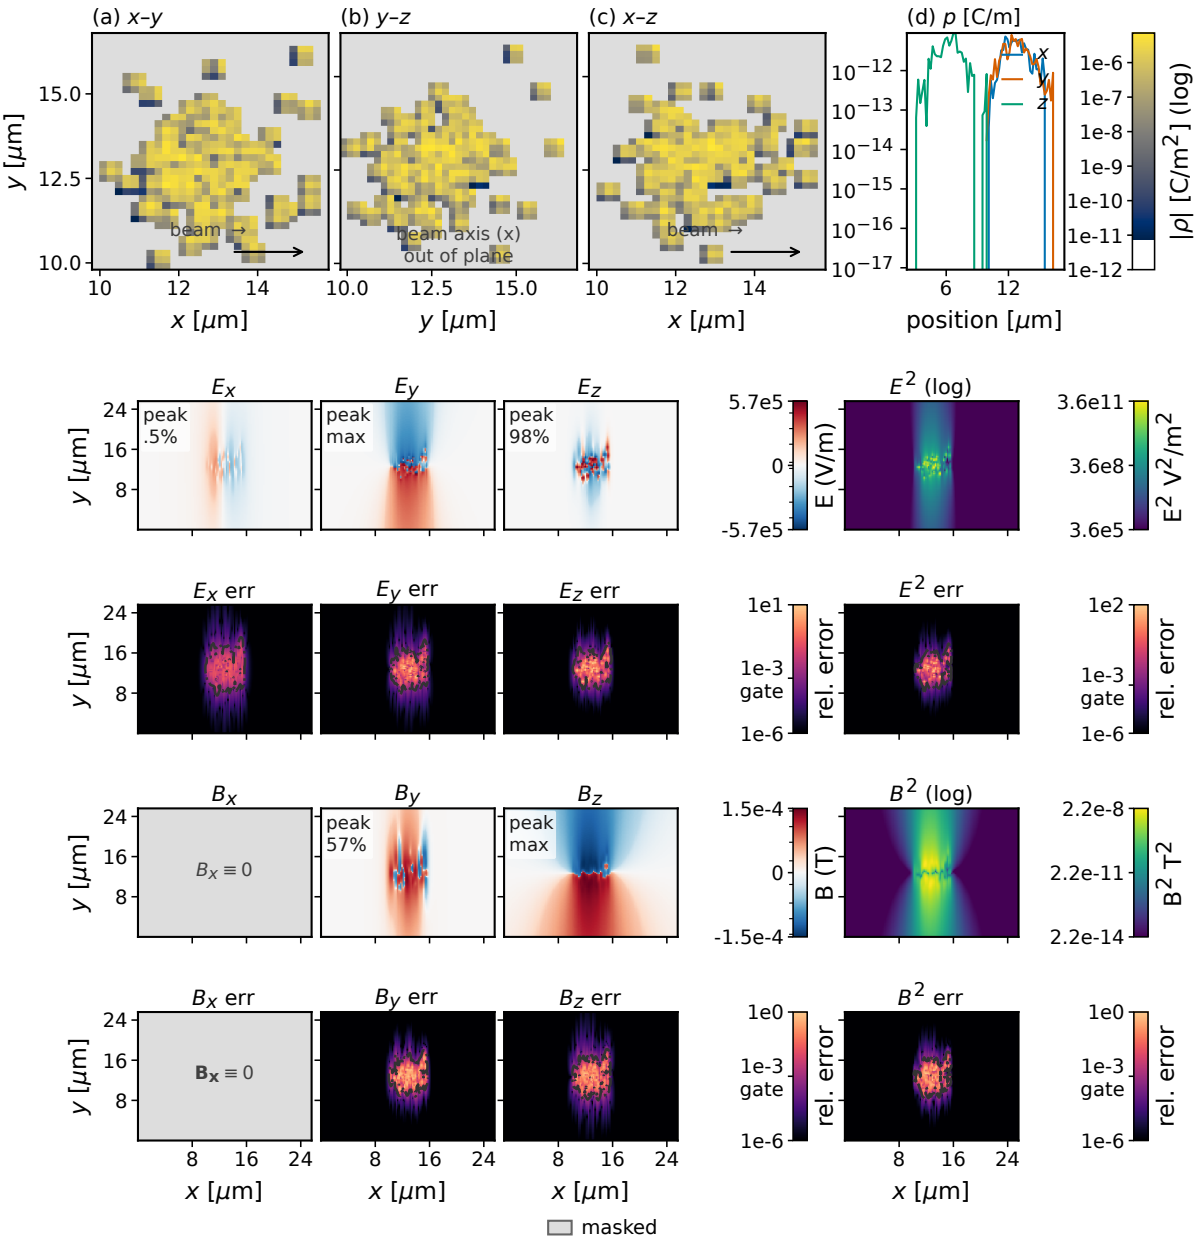


Figure 16: Large-domain slices ($\gamma = 10$, $128\times128\times64$, cell 2×10^{-7} m, $N = 100$, JURECA H100), axial-x geometry: the bunch moves along $+x$, so B_x, the component along the propagation axis, is identically zero. In the x-y cut the beam axis x is horizontal; the z-y cut is a transverse view with the beam axis x perpendicular to the page, so the horizontal axis is the transverse y, not the propagation direction, and z is vertical. The x-y plane is shown; the z-y plane is continued in Fig. 16. The top row is a compact charge-density overview of the same bunch: $|\rho|$ on a logarithmic scale floored at 10^{-6} of the peak, in the three coordinate planes on one shared colour bar, together with a marginal bunch profile. The `beam→` marker gives the propagation direction, and panels transverse to it are marked "beam axis (out of plane)". Below the density row, each image carries the electric block first and the magnetic block second. The electric block shows four columns E_x, E_y, E_z, E^2: the three components share one symmetric-logarithmic scale and bar (peak in percent of the shared maximum), E^2 has its own logarithmic scale and unit label, and a logarithmic error row against the order-4 driver reference lies beneath. E_x is suppressed but nonzero, peaking at 0.5475% (x-y) and 0.1384% (z-y) of the shared maximum. The magnetic block shows the same four-column layout and error convention for B_x, B_y, B_z, B^2. For uniform motion along $+x$ the field is purely transverse, so B_x is identically zero: its field panel is blank and its error panel is fully masked, since the relative error of an identically zero component is undefined. The magnetic plane peak is 1.5×10^{-4} T (x-y) and 1.4×10^{-4} T (z-y). Grey pixels are masked.

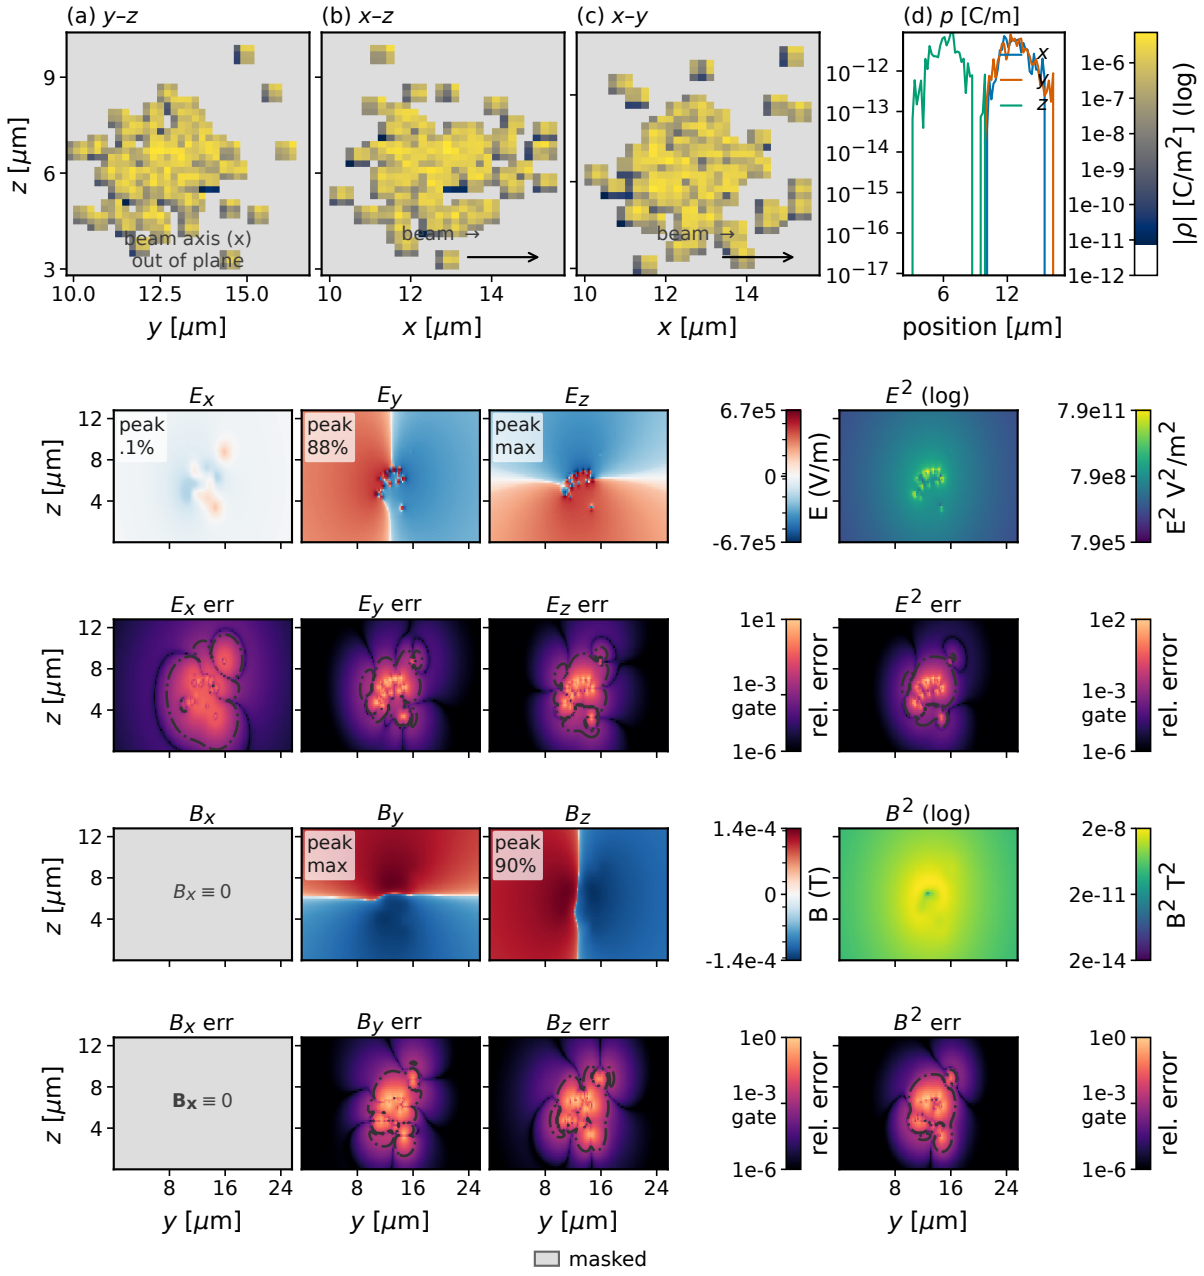


Figure 16: Large-domain slices (continued): z-y plane at $\gamma = 10$, with the same electric/magnetic blocks, colour scales, masks and error conventions as the x-y plane in Fig. 16.

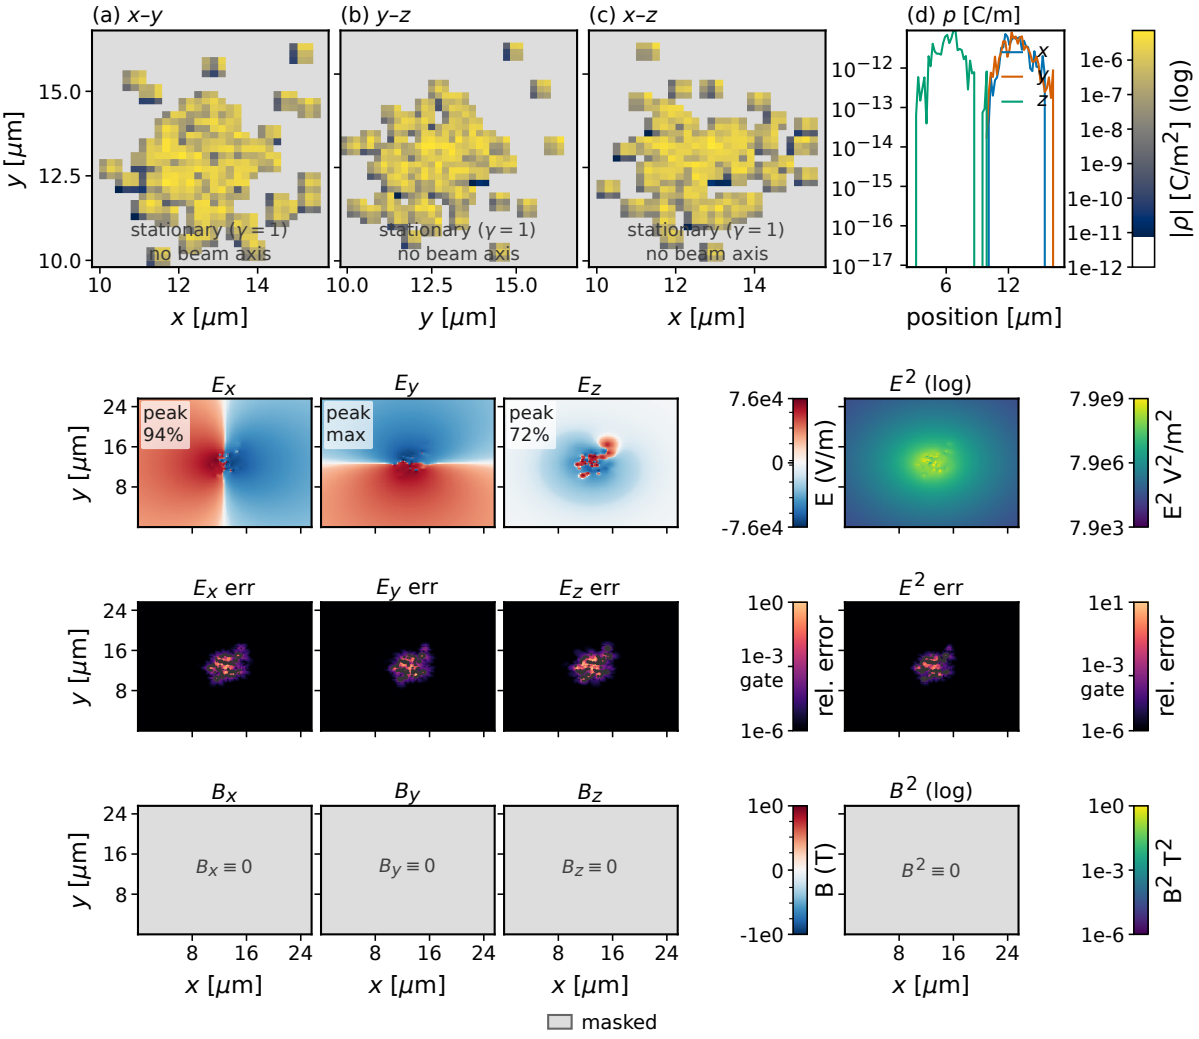


Figure 17: Large-domain slices for the stationary bunch ($\gamma = 1$), same domain and layout as Fig. 16. In the x-y cut x is horizontal; the z-y cut is the same transverse view as in Fig. 16, with x perpendicular to the page, y horizontal and z vertical; the bunch is stationary, so it has no propagation direction. The top row is a compact charge-density overview of the same bunch: $|\rho|$ on a logarithmic scale floored at 10^{-6} of the peak, in the three coordinate planes on one shared colour bar, together with a marginal bunch profile. The `beam`$\rightarrow$ marker gives the propagation direction, and panels transverse to it are marked "beam axis (out of plane)". The x-y plane is shown; the z-y plane is continued in Fig. 17. Below the density row, each image has the electric block first and the magnetic field row second. The electric block shows four columns E_x, E_y, E_z, E^2: the three components share one symmetric-logarithmic bar (peak in percent of the shared maximum), E^2 has its own logarithmic scale and unit label, and a logarithmic error row lies beneath. The core error reaches 8.4×10^{-1} (z-y E_x), far above the 10^{-3} gate. The magnetic field row shows four columns B_x, B_y, B_z, B^2; the stationary reference magnetic field is identically zero, so every magnetic field panel is blank. The magnetic error row is omitted because the relative error of an identically zero component is undefined, as in the per-γ $\gamma = 1$ figure. Grey pixels are masked.

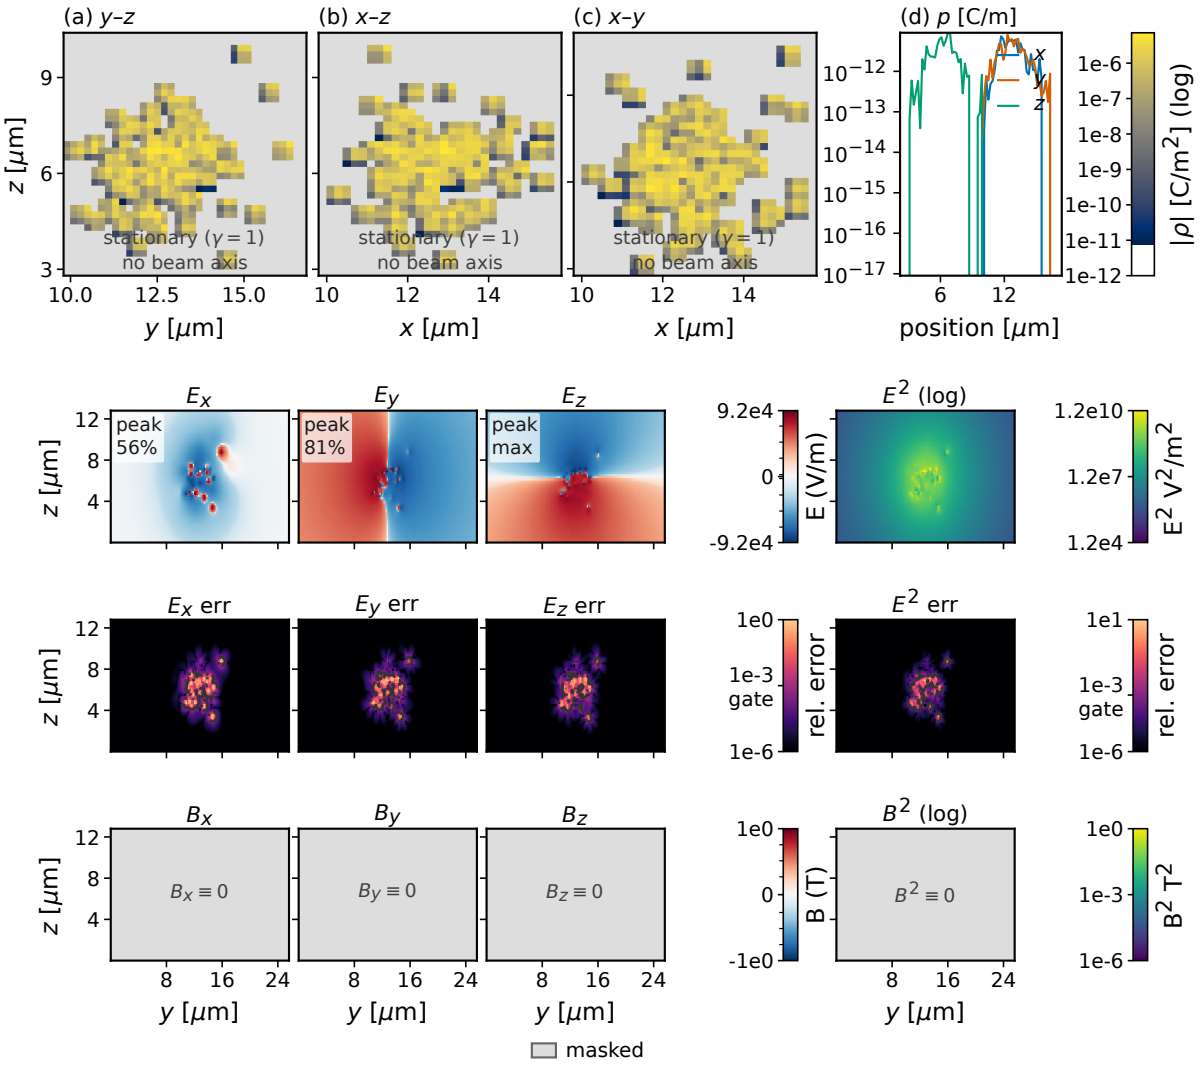


Figure 17: Large-domain slices for the stationary bunch (continued): z-y plane, with the same blocks, colour scales, masks and error conventions as the x-y plane in Fig. 17.

Table 17: Binned multi-velocity peaks, 64^3 grid of cell 10^{-7} m, 200 sources of weight 5×10^{-13} C ($Q = -100$ pC), $\gamma = 10$.

spread	B=16 $\lvert B_z \rvert$ (T)	B=64 $\lvert B_z \rvert$ (T)	B=16 $\lvert B_y \rvert$ (T)	B=64 $\lvert B_y \rvert$ (T)
energy 5%	2.32780×10^4	2.32778×10^4	0	0
angular 5°	2.48877×10^4	2.48872×10^4	1.2851×10^3	1.2864×10^3

D Multi-velocity binning, measured

D.1 Binned convolution

The single-velocity scope of Section 2 is extended by splitting the velocity distribution into B classes and running one convolution per class; the near shell needs no binning, because each pair carries its own velocity. The measurements below are on a 64^3 grid of cell 10^{-7} m with 200 sources of weight 5×10^{-13} C ($Q = -100$ pC) on JURECA H100 (H100 `sm_90`, CUDA 13.4).

Box representative. Each box is represented by its absolute-charge-weighted mean momentum, not by the geometric box centre. The mean minimises the angular and energy deviations for which the error model charges the split, at no extra kernel cost. On a 200-particle, $\gamma = 10$, 5° bunch the global L2 field error against the 128-box field falls from 13.2% to 11.2% at 16 boxes and from 4.31% to 4.13% at 64 boxes. The raw rows are in `docs/measurements/velocity_bin_representative.csv`.

Single-velocity sanity. With zero spread the binned path is bitwise equal to the unbinned path: $B = 1$ and $B = 16$ give bit-identical fields in all six components ($\max|\Delta| = 0$).

Figure 18 shows the sampled distributions and the bin comparison. With a 5% energy spread and $+y$ drift the B_z peak is 2.32780×10^4 T at 16 bins and 2.32778×10^4 T at 64 bins, 6.9×10^{-6} relative apart; $B_y = 0$ exactly, which is correct for a purely longitudinal current. With a 5° angular spread the B_z peak moves from 2.48877×10^4 to 2.48872×10^4 T and B_y from 1285.1 to 1286.4 T. The B_y peak is the small transverse component, 1.29×10^3 T against B_z 2.49×10^4 T, so its relative deviation, 0.107%, is the sensitive binning metric; the B_z peaks agree to 2.33×10^{-5}. The ratio B_y/B_z is 5.16×10^{-2}. These are fresh measurements of the committed bunch on the final code; they supersede the earlier tabulated angular peaks from code `c0bf82c`.

1 mrad case against the arbiter. The fluctuation bunch (Table 13 row "1 mrad, binned": $N = 200$, weight 5×10^{-13} C each, $Q = -100$ pC, $\gamma = 10$, $+y$ drift, 1.046 mrad rms x-divergence) is measured at 64 bins on the 64^3 grid of cell 10^{-7} m and compared against the Gauss–Legendre arbiter (unbinned, per-particle u, 698-node slice subset) as the reference. On the x-y slice at $z = 32$ cells ($z = 3.2$ µm, the grid centre) picBunch gives $B_y = 7.835$ T and $B_z = 2.3336 \times 10^4$ T, against arbiter values of 7.759 T and 2.344×10^4 T (factors 1.010× and 0.996×); on the z-y slice at $x = 32$ cells ($x = 3.2$ µm) the corresponding values are 7.561 T / 7.519 T and 2.3159×10^4 T / 2.327×10^4 T (1.006× and 0.995×). The 3-D B_y peak is 8.734 T and B_z is 2.3336×10^4 T. The comparison criterion is the per-component slice peaks—peak $|B_y|$ and peak $|B_z|$ on each plane—and the B_y/B_z ratio; the worst deviation is 1.4% (x-y ratio) and every per-component peak on every slice is within 1.0%. The 1 mrad case was re-measured from the archived input on the reduced $32 \times 32 \times 16$ slice: it reproduces the committed field to machine precision (worst relative peak difference 7.1×10^{-16}). The campaign reference differs from `indfield.py` within two cells of a source by at most 9.6×10^{-4} of the E peak outside that region. The full 64^3 peaks above and the 1.4% criterion are not re-measured here; they keep their original provenance, the full 64^3 JURECA H100 run.

Runtime: $B = 64$, 64^3, 200 particles, `-near-shell fast` ($R = 3$, tile 4, order 5), 248 s on an A30.

Far-tail switch validation

For the documented 20-cell, 4096-electron bunch (source radius $R = 16.85$ cells, peak 45722.5 V/m, tolerance level 45.72 V/m), the quadrupole margin of Eq. (12) is 350 cells, which exceeds the 2^{28} padded-node cap, so the code falls back to the measured far-tail switch. Table 18 gives the switch error: the worst per-component $|\mathbf{E}_{\text{tail}} - \mathbf{E}_{\text{direct}}|$ against an independent direct charge sum in 13 directions.

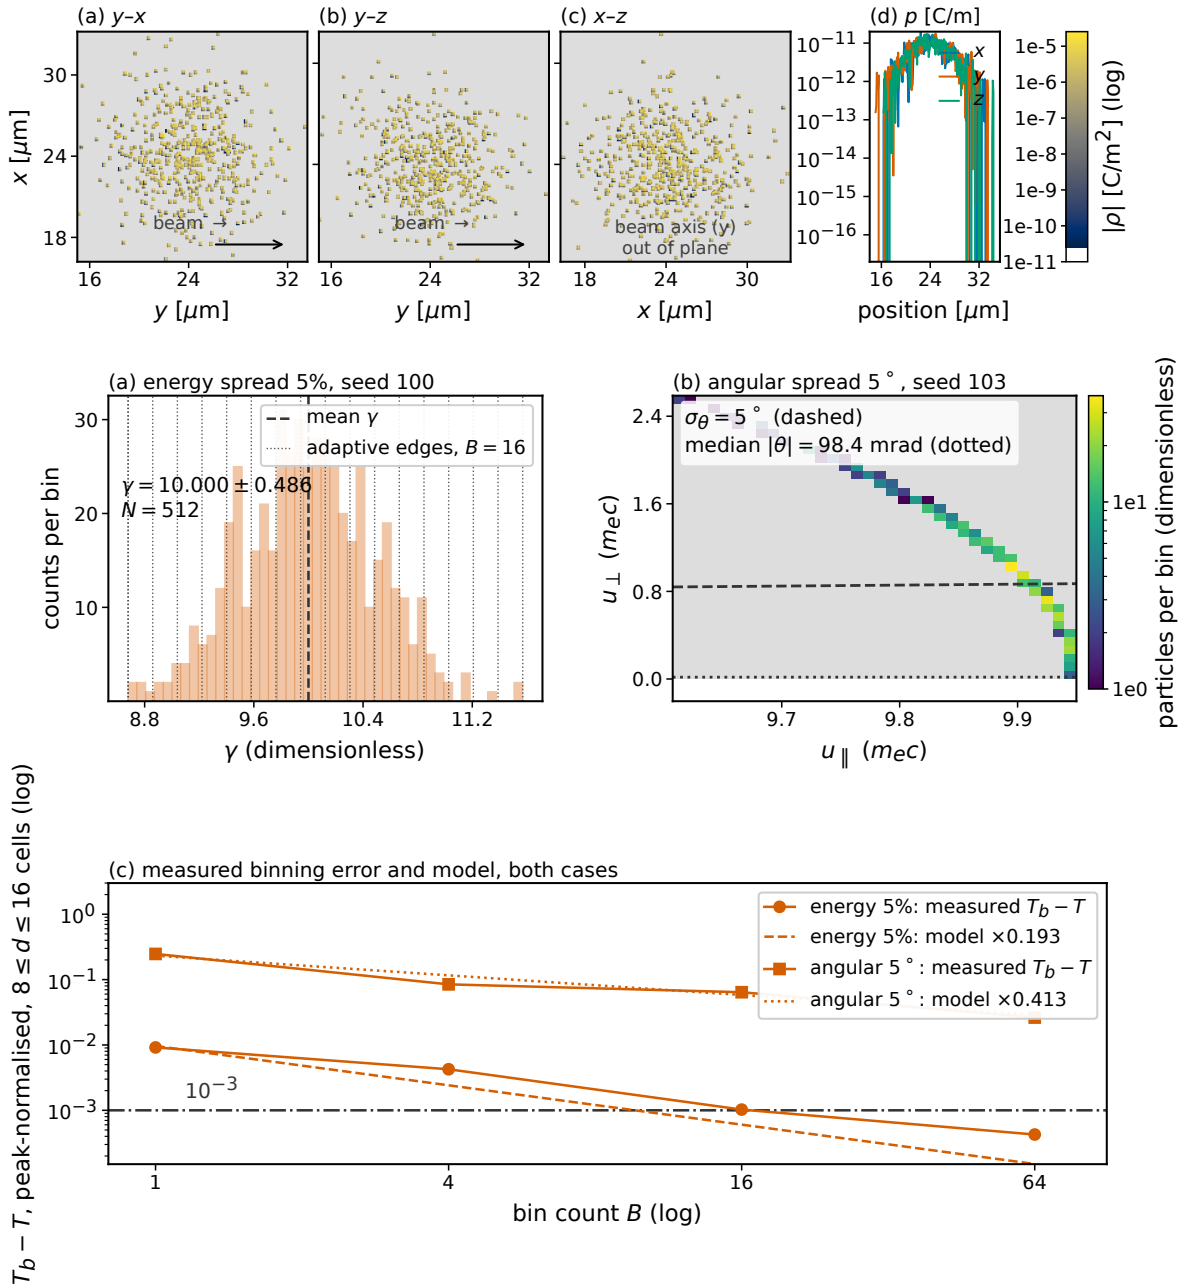


Figure 18: Sampled multi-velocity bunches ($\gamma = 10$, 512 sources, seeds 100 and 103); the $B = 16$ edges in (a) and the velocity boxes that produce the field come from the same `VelocityBinning` per-axis min/max path. (a) γ histogram for the 5% energy spread with the 16 adaptive edges drawn and the mean and standard deviation annotated; (b) $u_{\parallel}$–$u_{\perp}$ density for the 5° angular spread with the nominal 5° divergence and the measured median $|\theta| = 101.8$ mrad marked; (c) relative deviation of the $B = 16$ peaks from the $B = 64$ reference, 6.9×10^{-6} (energy B_z), 2.3×10^{-5} (angular B_z) and 1.1×10^{-3} (0.107%, angular B_y), with $B_y \equiv 0$ for the energy case, so the field is bin-independent at the stated $2.32780 \times 10^4 / 2.32778 \times 10^4$ T and 5.16×10^{-2} ratio. The B_y peak is 1.29×10^3 T against B_z 2.49×10^4 T, so the relative B_y deviation is the sensitive binning metric. The top row is a compact charge-density overview of the same bunch: $|\rho|$ on a logarithmic scale floored at 10^{-6} of the peak, in the three coordinate planes on one shared colour bar, together with a marginal bunch profile. The `beam`→ marker gives the propagation direction, and panels transverse to it are marked "beam axis (out of plane)".

A 512^3 output sampled outside the compute box confirms the switch in practice: 36 nodes including corners, worst $|\mathbf{E}_{\text{file}} - \mathbf{E}_{\text{direct}}| = 2.8 \times 10^{-3}$ V/m against the 45.7 V/m tolerance level (6.1×10^{-5} of it), and every sample nonzero, i.e. no interior zero fill.

Table 18: Far-tail switch validation, 20-cell 4096-electron bunch, $R = 16.85$ cells, tolerance level 45.72 V/m (10^{-3} of the peak 45722.5 V/m). Worst per-component $|\mathbf{E}_{\text{tail}} - \mathbf{E}_{\text{direct}}|$ against an independent direct charge sum over 13 directions. The switch at $6R$ is the chosen switch: the error is 1.87×10^{-3} of the tolerance level. rosi_v100 run: the direct-charge-sum comparison data are not in this repository, so these values are not re-measured here.

switch radius	worst err (V/m)	/ tolerance level
$2R$ (33.7 cells)	4.80×10^{1}	1.05
$4R$ (67.4 cells)	8.77×10^{-1}	1.92×10^{-2}
$6R$ (101.1 cells)	8.53×10^{-2}	1.87×10^{-3}
$8R$ (134.8 cells)	1.66×10^{-2}	3.63×10^{-4}

References

E. Esarey, C. B. Schroeder, and W. P. Leemans. Physics of laser-driven plasma-based electron accelerators. *Reviews of Modern Physics*, 81(3):1229–1285, 2009. doi:10.1103/RevModPhys.81.1229.

John David Jackson. *Classical Electrodynamics*. Wiley, Hoboken, NJ, 3 edition, 2007. ISBN 978-0-471-30932-1.

Lev Davidovich Landau and Evgeny Mikhailovich Lifshitz. *The Classical Theory of Fields*. Pergamon Press, Oxford, 2 edition, 1975. ISBN 978-0-08-025072-4.

C. K. Birdsall and A. B. Langdon. *Plasma Physics via Computer Simulation*. Wiley, New York, 1985. ISBN 0-070-05371-5.

R. W. Hockney and J. W. Eastwood. *Computer Simulation Using Particles*. Wiley, New York, 1981.

A. Burau, Réne Widera, J. Honig, J. Juckeland, A. Debus, T. Kluge, U. Schramm, M. D. Cowan, R. Sauerbrey, and Michael Bussmann. Picongpu: A fully relativistic particle-in-cell code for a gpu cluster. *IEEE Transactions on Plasma Science*, 38(10), 2010. doi:10.1109/tps.2010.2064310.

Axel Huebl, Rémi Lehe, Jean-Luc Vay, David P. Grote, Ivo Sbalzarini, Stephan Kuschel, David Sagan, Frédéric Pérez, Fabian Koller, and Michael Bussmann. openpmd 1.1.0: Base paths for mesh- and particle-only files and updated attributes, 2018. doi:10.5281/zenodo.1167843.

W. H. Press, S. A. Teukolsky, W. T. Vetterling, and B. P. Flannery. *Numerical Recipes: The Art of Scientific Computing*. Cambridge University Press, Cambridge, 1986. ISBN 978-0-521-30811-9.

NVIDIA Corporation. CUDA toolkit documentation: cufft library user guide. NVIDIA Corporation, Santa Clara, CA, 2024. URL https://docs.nvidia.com/cuda/cuda-cufft-api/index.html. Accessed 2026.

picBunch project. picbunch checkpoint, 4-rank PIConGPU restart on a decomposable grid (24 x 24 x 96, -d 1 1 4). docs/measurements/picongpu_restart_4rank.md, with archived checkpoints, logs and configuration, in the repository, 2026.